

Great Ideas in Computer Architecture

Pipelining and Hazards

Instructor: Steven Ho



Great Idea #4: Parallelism

Software

- **Parallel Requests**
Assigned to computer
e.g. search “Garcia”

- **Parallel Threads**
Assigned to core
e.g. lookup, ads

- **Parallel Instructions**
> 1 instruction @ one time
e.g. 5 pipelined instructions

- **Parallel Data**
> 1 data item @ one time
e.g. add of 4 pairs of words

- **Hardware descriptions**
All gates functioning in parallel at same time

Hardware

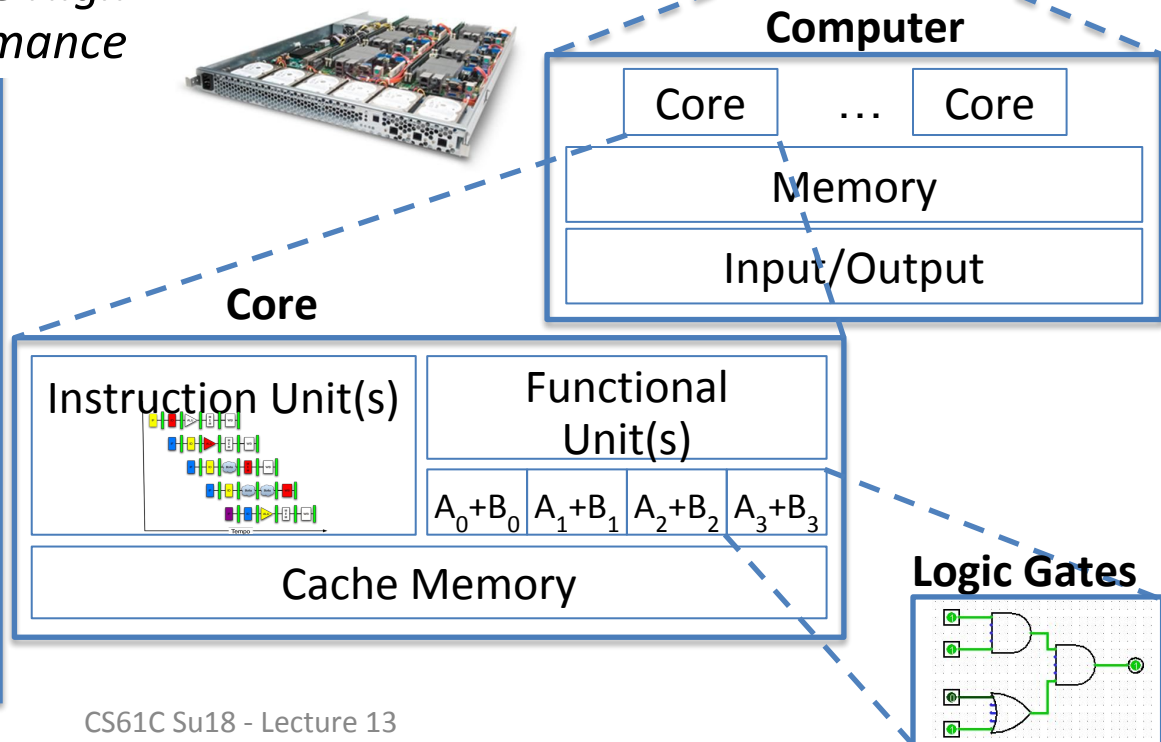
Warehouse
Scale
Computer



Smart
Phone



*Leverage
Parallelism &
Achieve High
Performance*



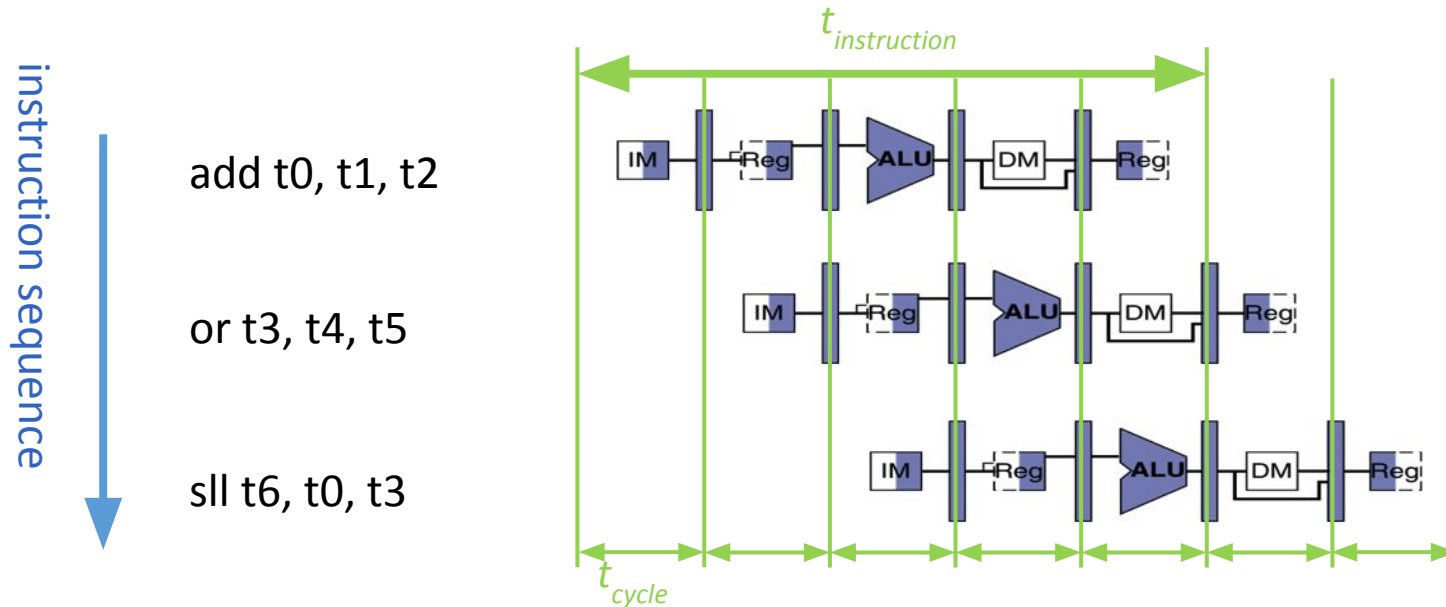
Review of Last Lecture

- Implementing controller for your datapath
 - Take decoded signals from instruction and generate control signals
- Pipelining improves performance by exploiting Instruction Level Parallelism
 - 5-stage pipeline for RISC-V: IF, ID, EX, MEM, WB
 - Executes multiple instructions in parallel
 - Each instruction has the same latency
 - What can go wrong???

Agenda

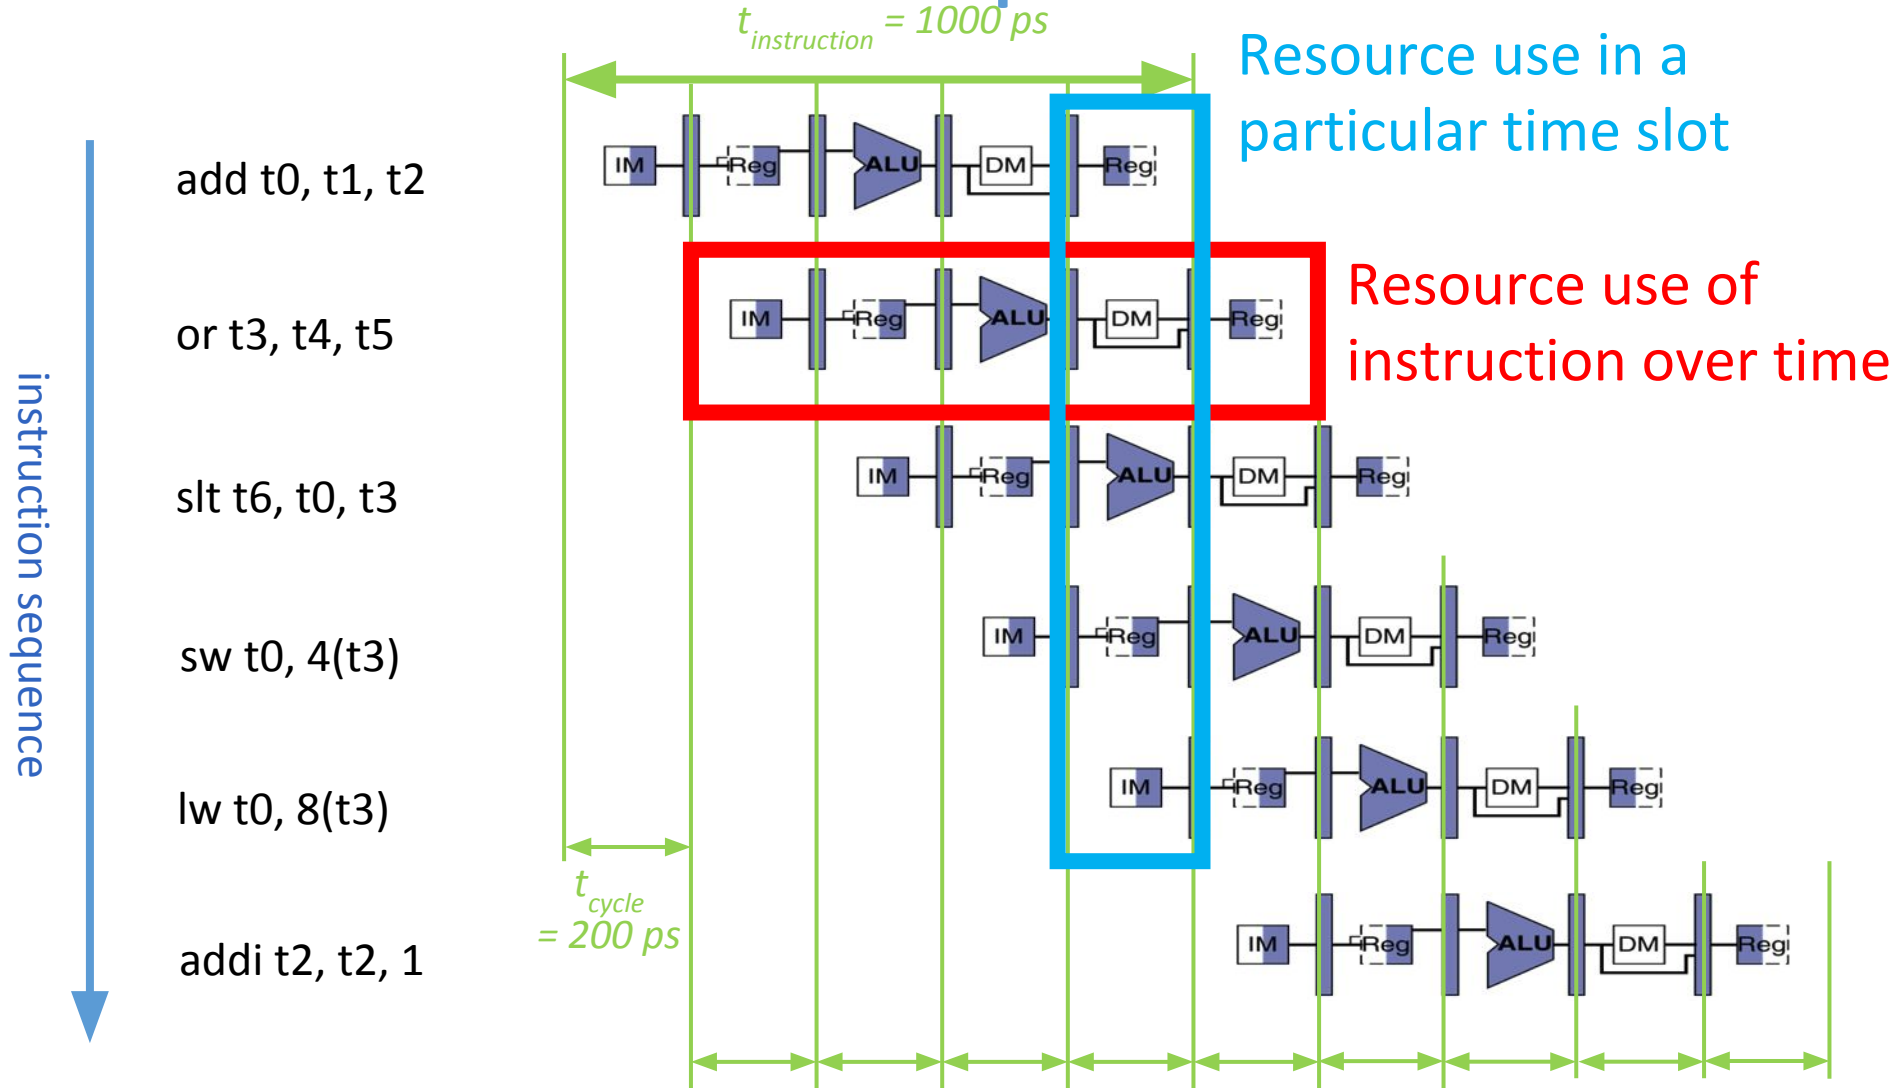
- RISC-V Pipeline
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- Superscalar processors

Recap: Pipelining with RISC-V

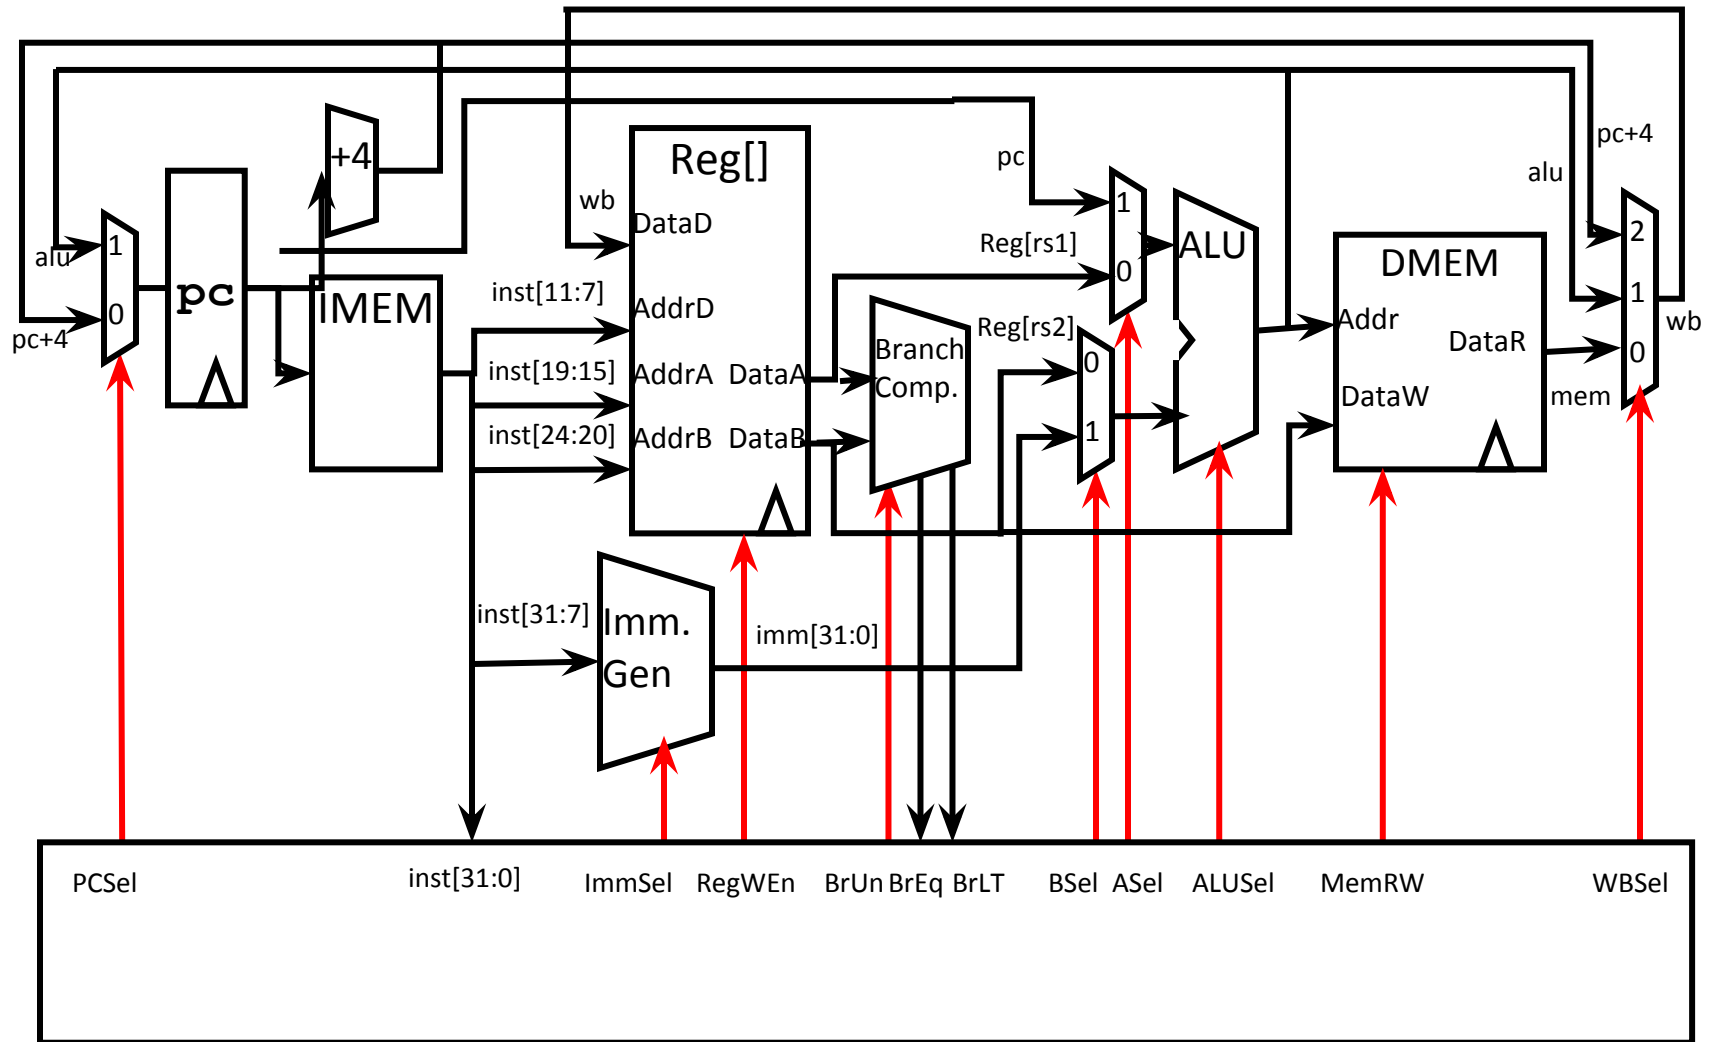


	Single Cycle	Pipelining
Timing	$t_{step} = 100 \dots 200 \text{ ps}$	$t_{cycle} = 200 \text{ ps}$
	Register access only 100 ps	All cycles same length
Instruction time, $t_{instruction}$	$= t_{cycle} = 800 \text{ ps}$	1000 ps
Clock rate, f_s	$1/800 \text{ ps} = 1.25 \text{ GHz}$	$1/200 \text{ ps} = \text{5 GHz}$
Relative speed	1 x	4 x

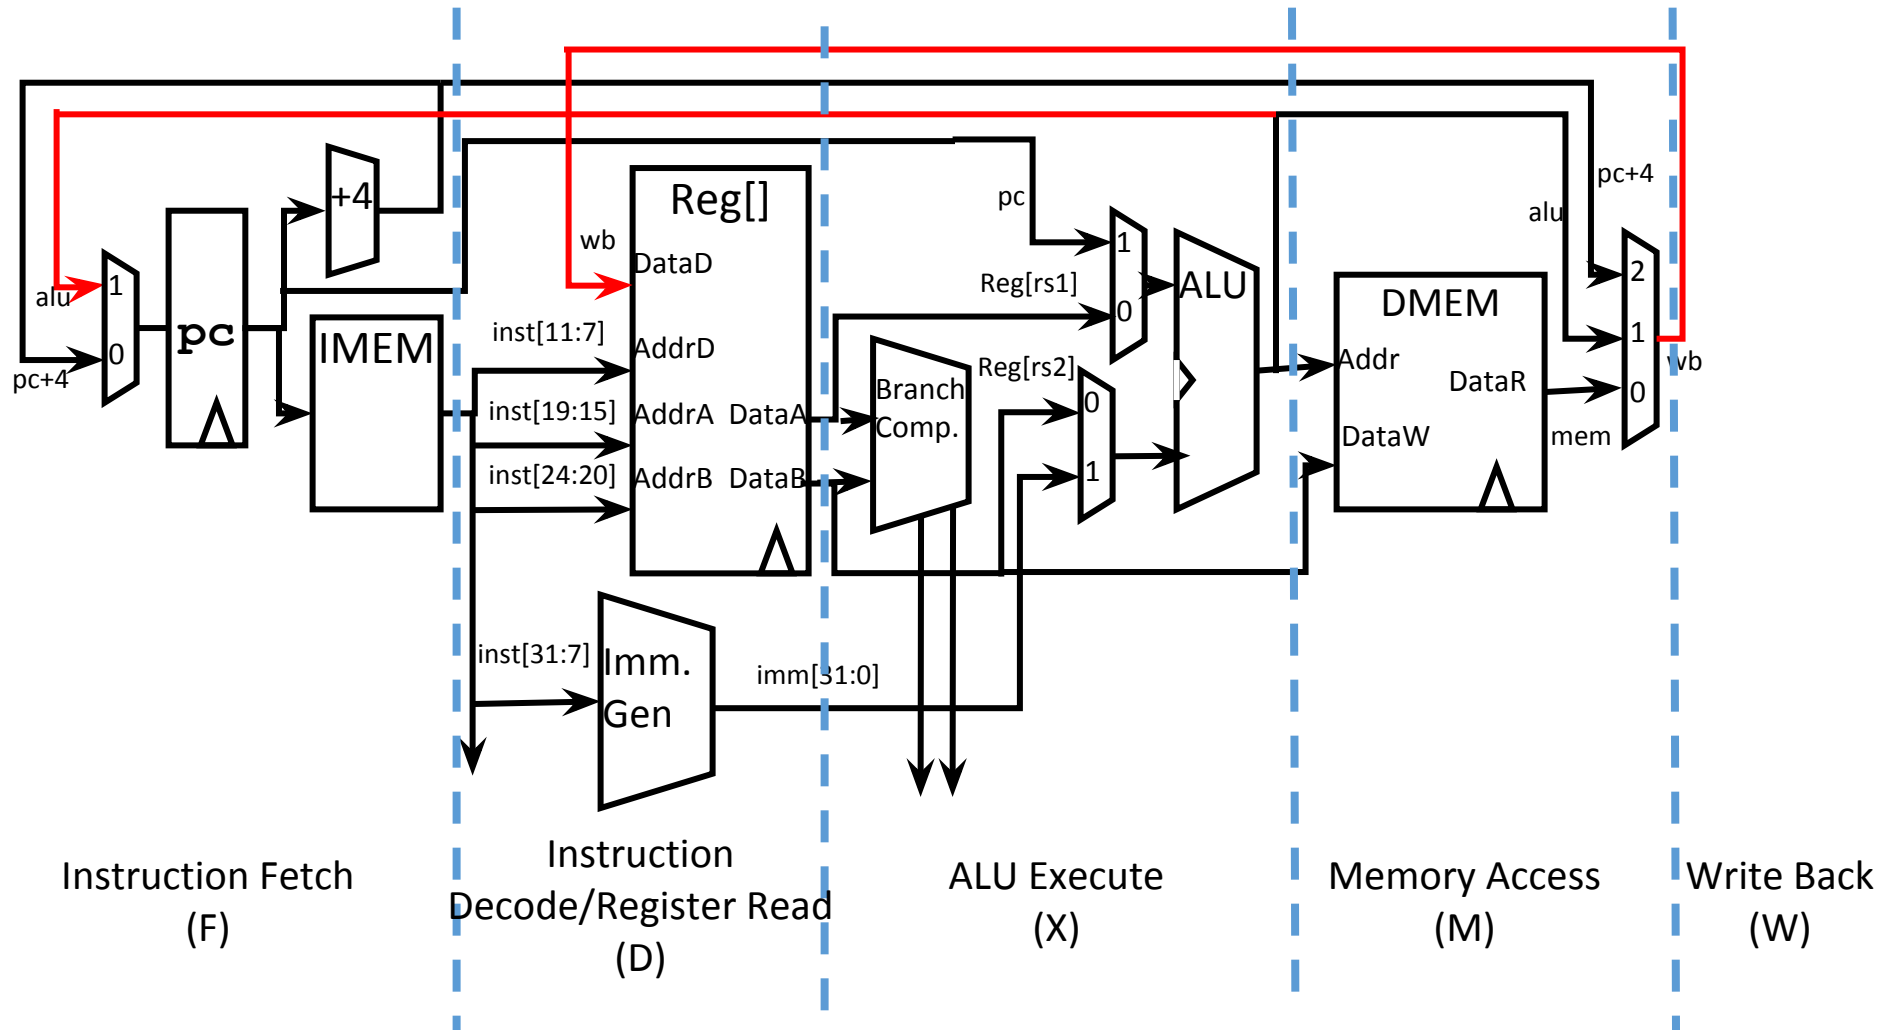
RISC-V Pipeline



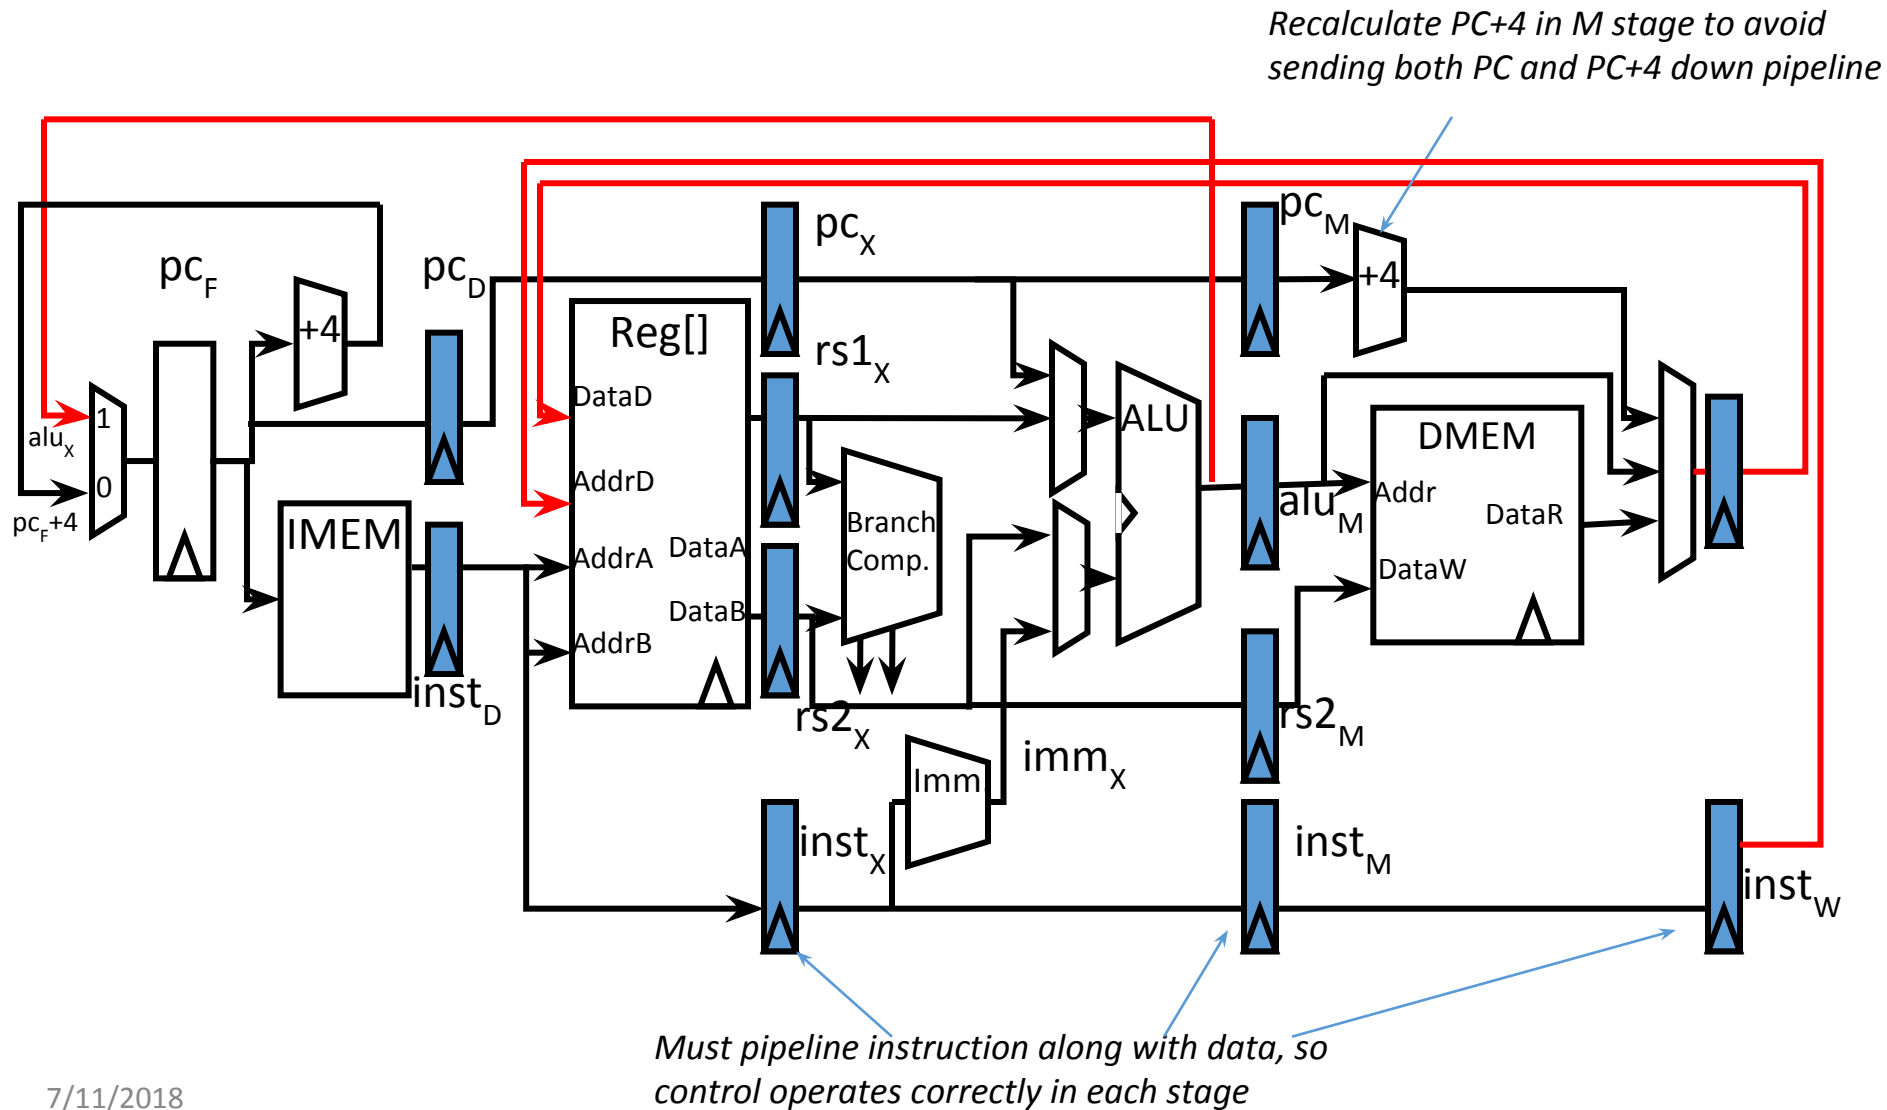
Single-Cycle RISC-V RV32I Datapath



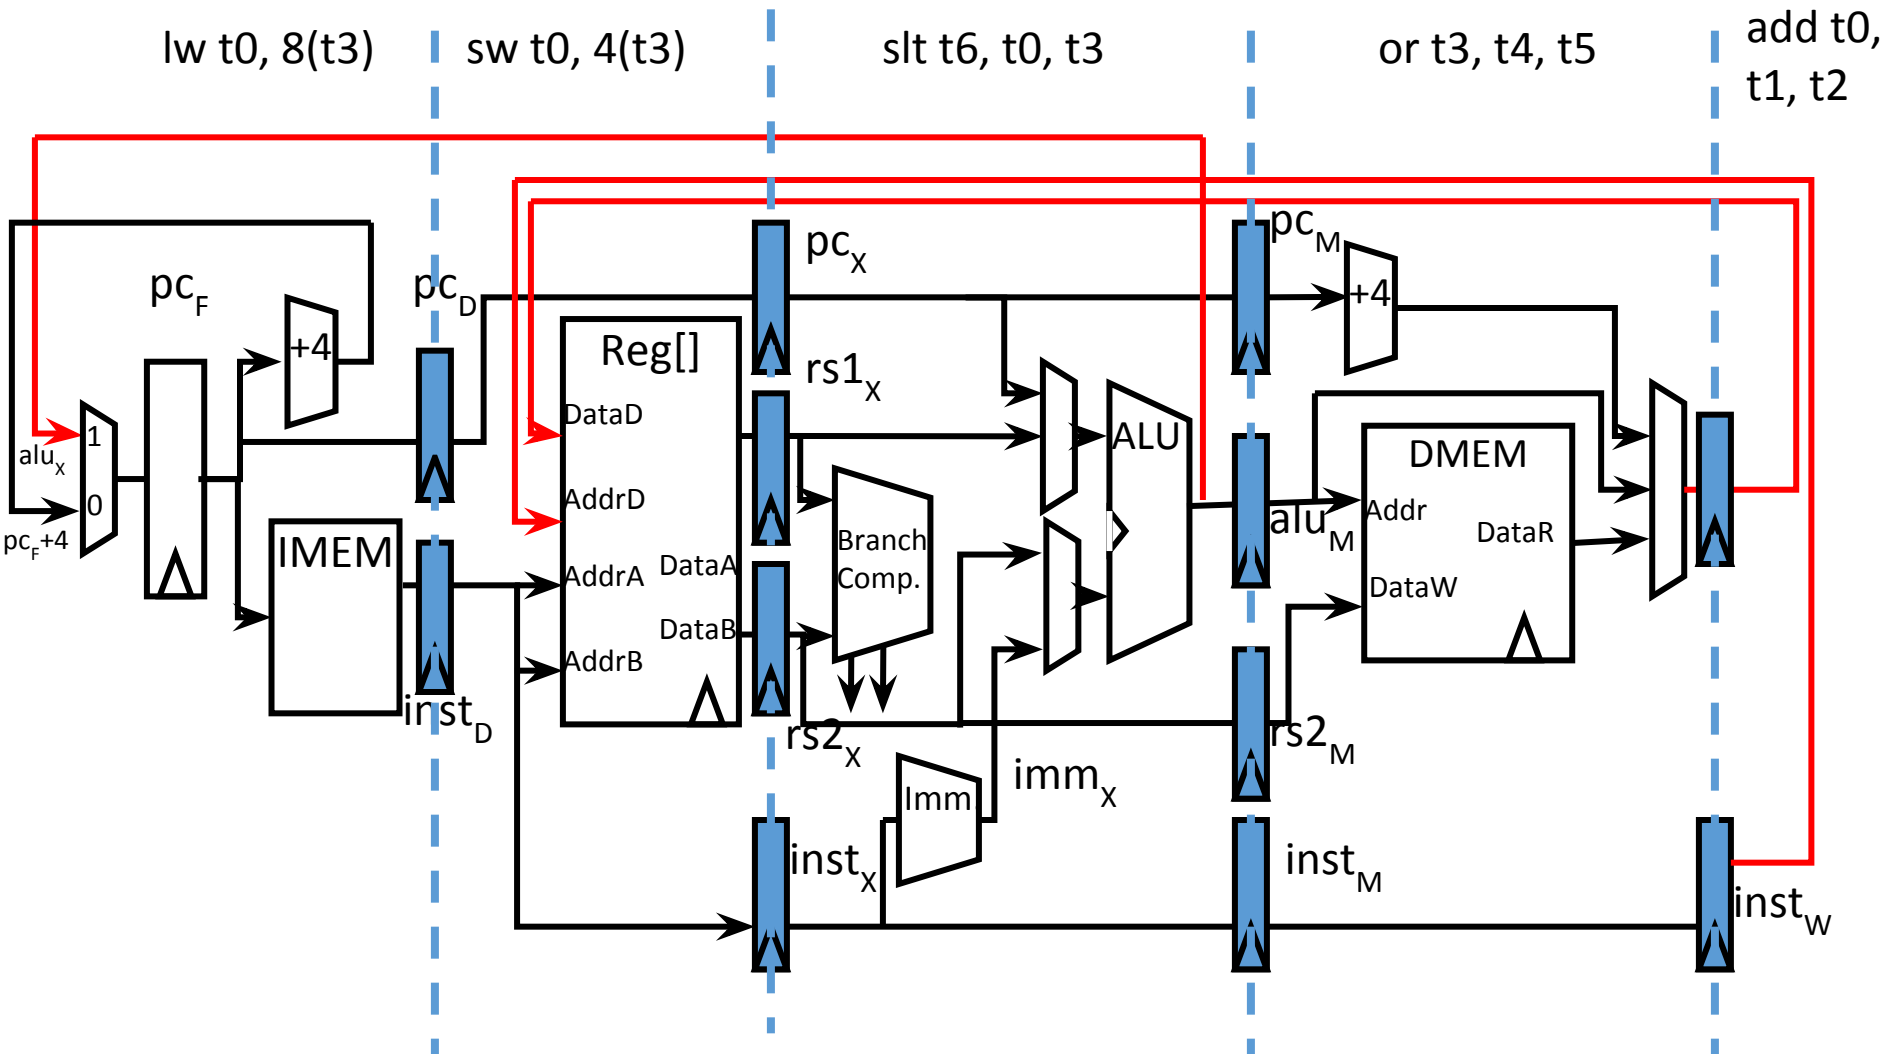
Pipelining RISC-V RV32I Datapath



Pipelined RISC-V RV32I Datapath




Each stage operates on different instruction



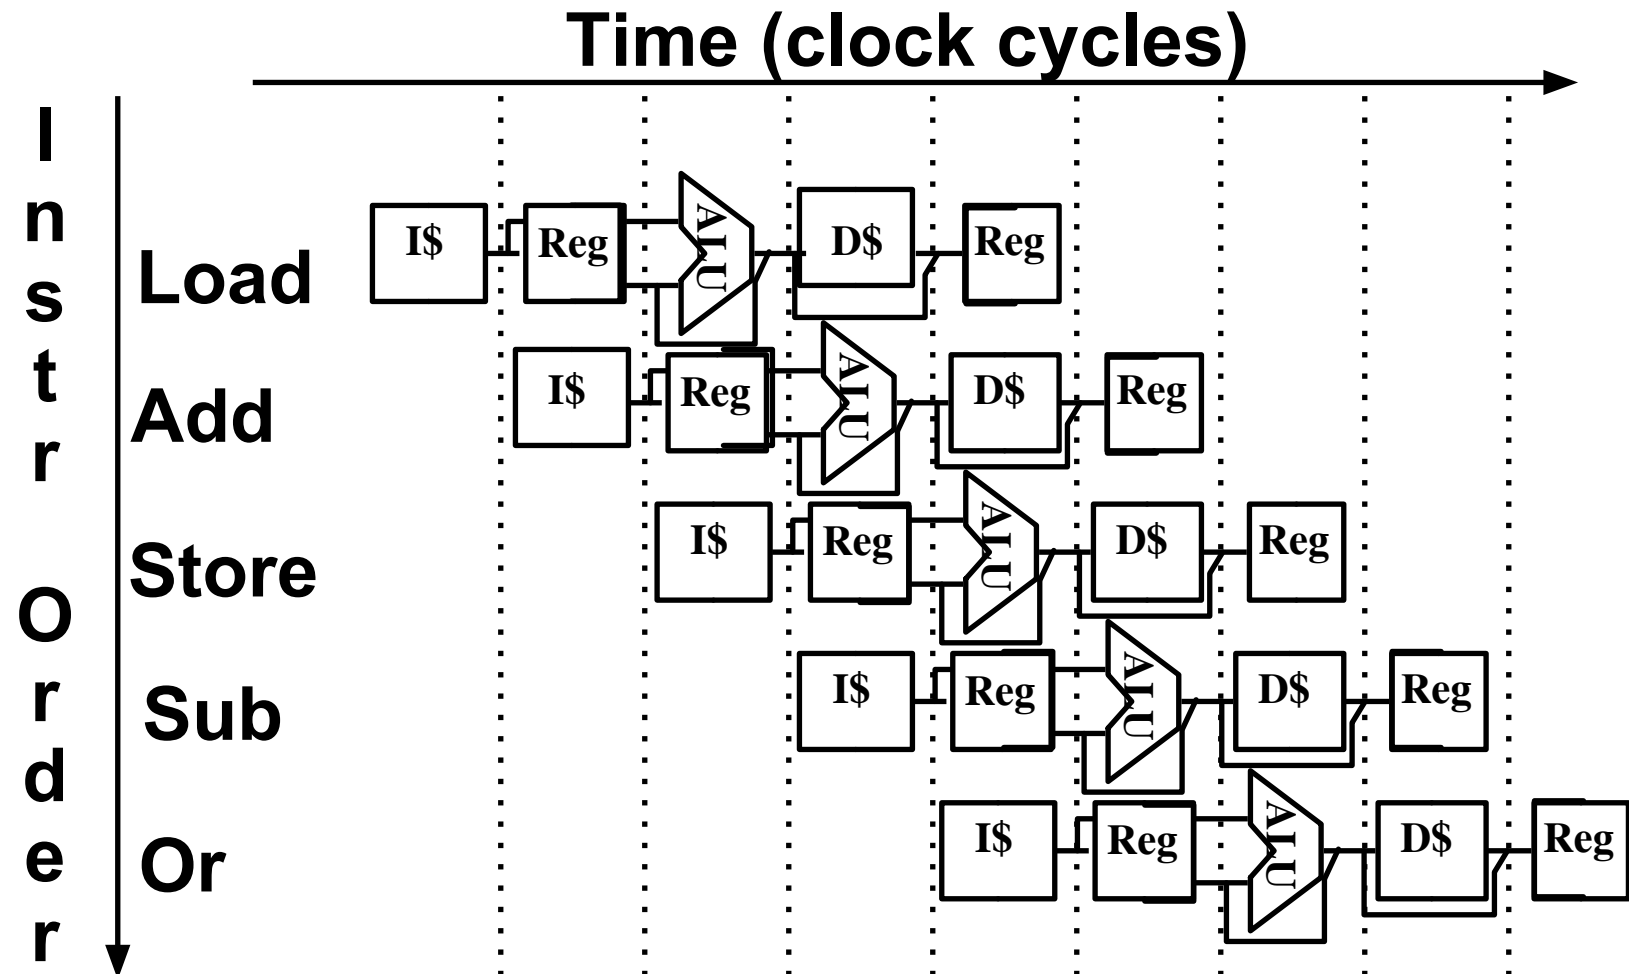
Pipelined Control

- Control signals derived from instruction
 - As in single-cycle implementation
 - Information is stored in pipeline registers for use by later stages



Graphical Pipeline Representation

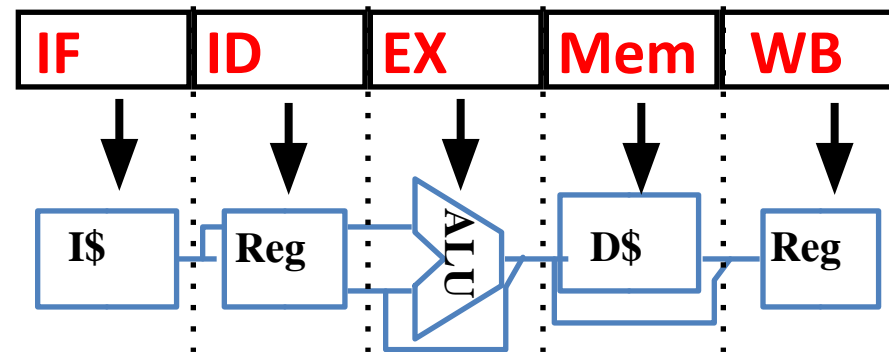
- RegFile: right half is read, left half is write



Question: Which of the following signals (buses or control signals) for RISC-V does NOT need to be passed into the EX pipeline stage for a beq instruction?

beq t0 t1 Label

- (A) BrUn
- (B) MemWr
- (C) RegWr
- (D) WBSeI



Agenda

Hazards Ahead!

- RISC-V Pipeline
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- Superscalar processors



Pipelining Hazards

A *hazard* is a situation that prevents starting the next instruction in the next clock cycle

1) *Structural hazard*

- A required resource is busy
(e.g. needed in multiple stages)

2) *Data hazard*

- Data dependency between instructions
- Need to wait for previous instruction to complete its data write

3) *Control hazard*

- Flow of execution depends on previous instruction

Agenda

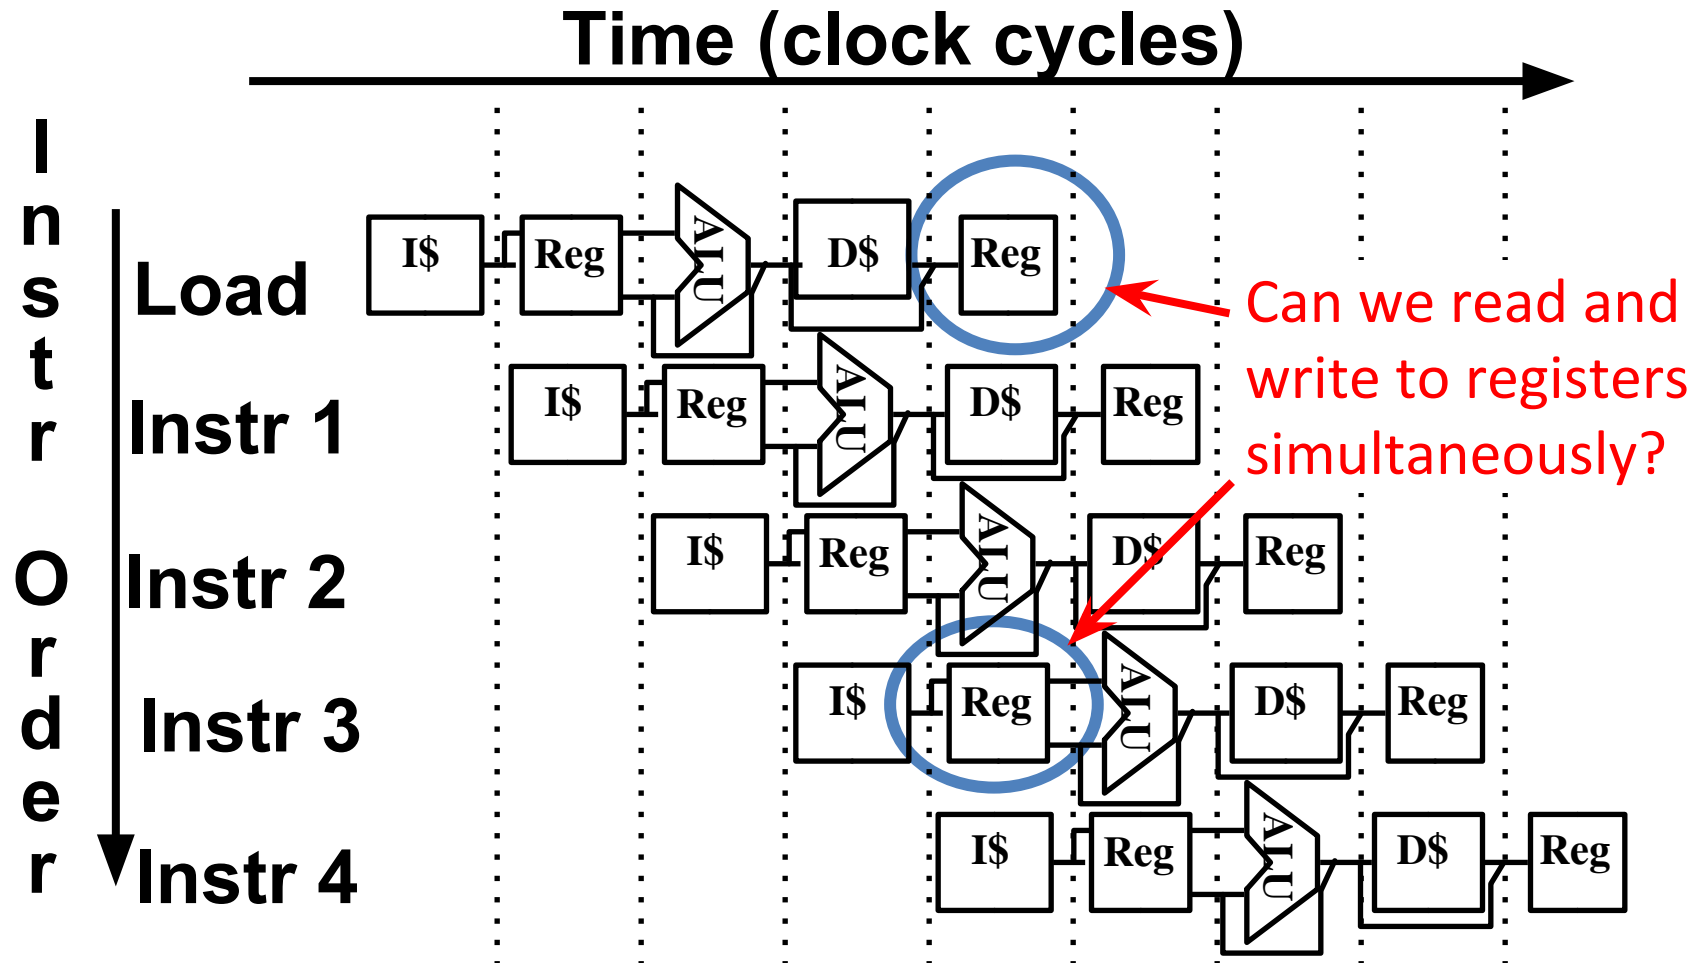
- RISC-V Pipeline
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- Superscalar processors

Structural Hazard

- **Problem:** Two or more instructions in the pipeline compete for access to a single physical resource
- **Solution 1:** Instructions take it in turns to use resource, some instructions have to stall
- **Solution 2:** Add more hardware to machine
- Can always solve a structural hazard by adding more hardware

1. Structural Hazards

- Conflict for use of a resource



Regfile Structural Hazards


- Each instruction:
 - can read up to two operands in decode stage
 - can write one value in writeback stage
- Avoid structural hazard by having separate “ports”
 - two independent read ports and one independent write port
- Three accesses per cycle can happen simultaneously

Regfile Structural Hazards

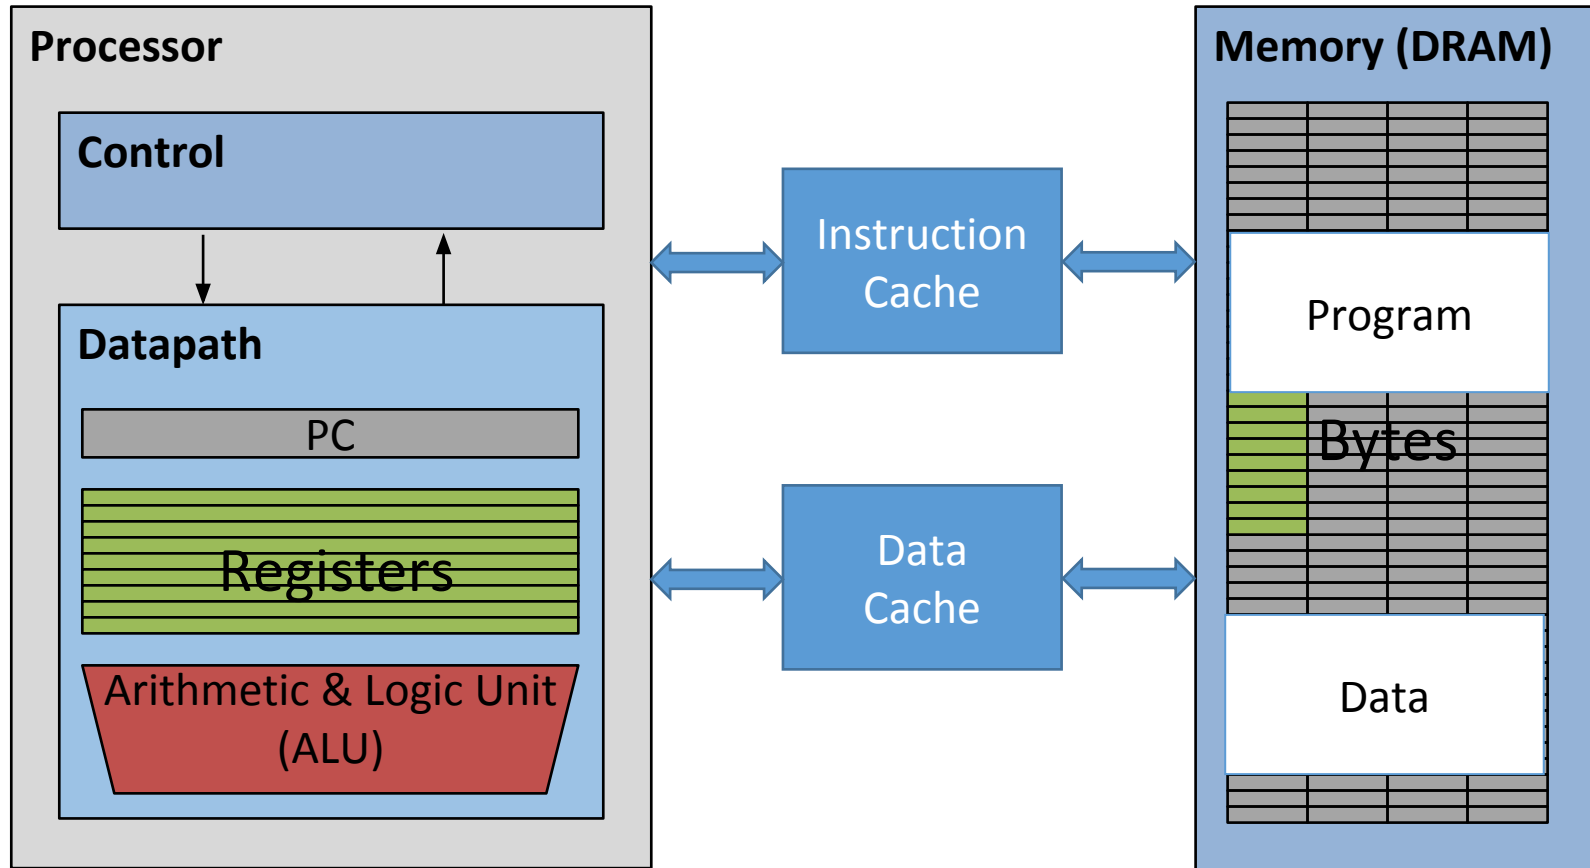
- Two *alternate* solutions:
 - 1) Build RegFile with independent read and write ports (what you will do in the project; good for single-stage)
 - 2) Double Pumping: split RegFile access in two! Prepare to write during 1st half, write on falling edge, read during 2nd half of each clock cycle
 - Will save us a cycle later...
 - Possible because RegFile access is *VERY* fast (takes less than half the time of ALU stage)
- **Conclusion: Read and Write to registers during same clock cycle is okay**

Memory Structural Hazards

- Conflict for use of a resource



Instruction and Data Caches



Caches: small and fast “buffer” memories

Structural Hazards – Summary

- Conflict for use of a resource
- In RISC-V pipeline with a single memory
 - Load/store requires data access
 - Without separate memories, instruction fetch would have to *stall* for that cycle
 - All other operations in pipeline would have to wait
- Pipelined datapaths require separate instruction/data memories
 - Or separate instruction/data caches
- RISC ISAs (including RISC-V) designed to avoid structural hazards
 - e.g. at most one memory access/instruction

Administrivia

- Proj2-2 due 7/13, HW3/4 due 7/16
 - 2-2 autograder being run
- Guerilla Session Tonight! 4-6pm
- HW0-2 grades should now be accurate on glookup
- Project 3 released tomorrow night!
- Supplementary review sessions starting
 - First one this Sat. (7/14) 12-2p, Cory 540AB

Agenda

- RISC-V Pipeline
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- Superscalar processors

2. Data Hazards (1/2)

- Consider the following sequence of instructions:

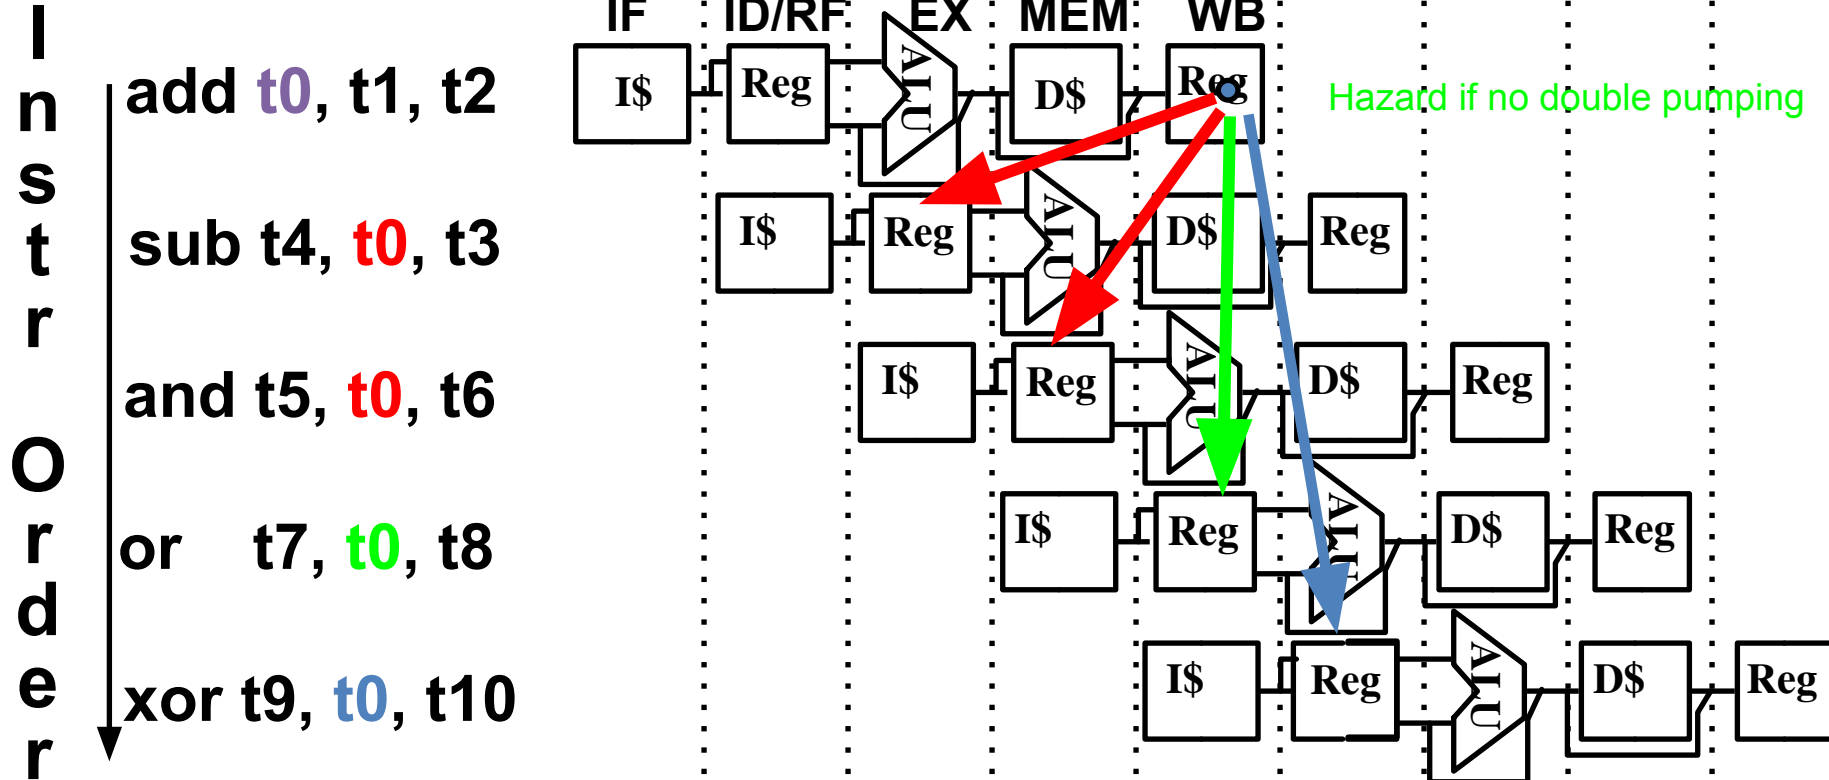
```
add  t0, t1, t2
sub  t4, t0, t3
and  t5, t0, t6
or   t7, t0, t8
xor  t9, t0, t10
```

Stored during WB Read during ID

2. Data Hazards (2/2)

- Data-flow *backward* in time are hazards

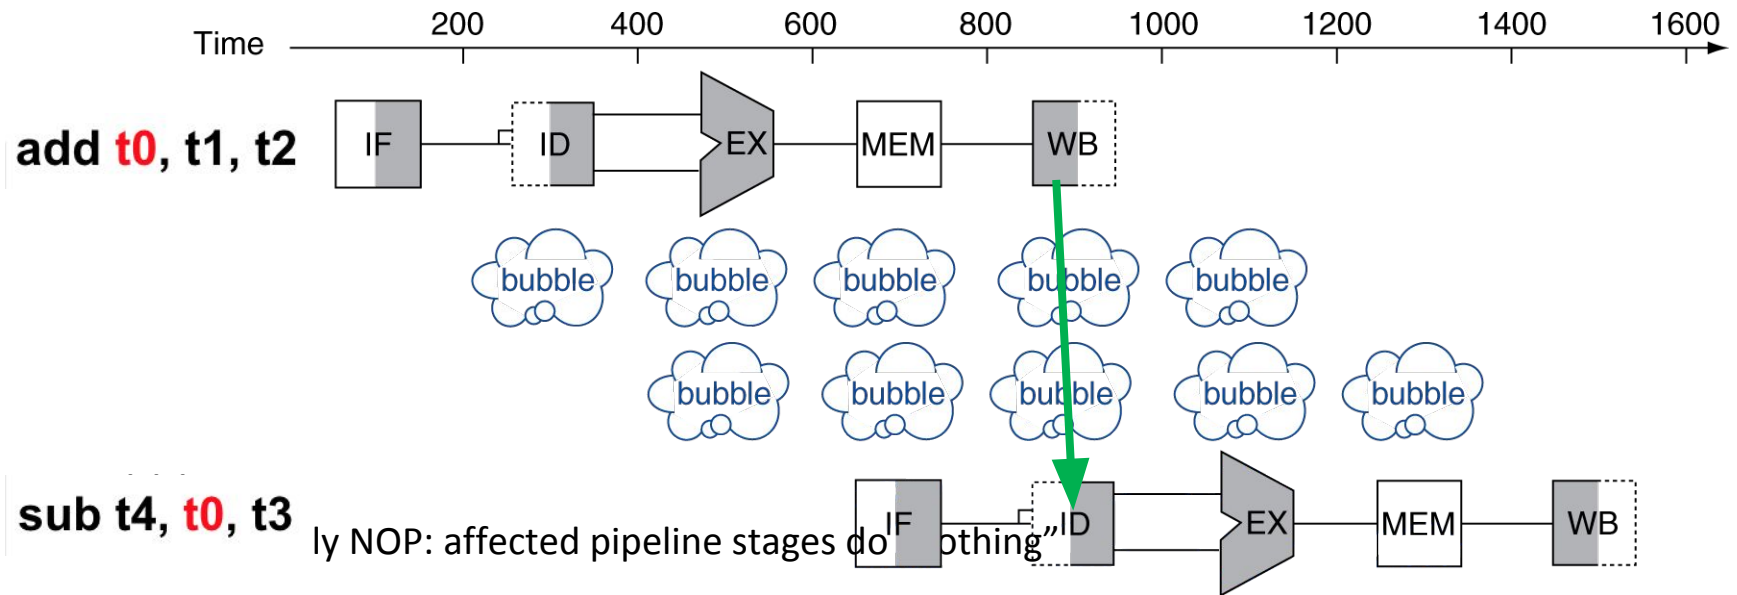
Time (clock cycles)



Solution 1: Stalling

- Problem: Instruction depends on result from previous instruction

- add t0, t1, t2
sub t4, t0, t3



Stalls and Performance

- Stalls reduce performance
 - But stalls are required to get correct results
- Compiler can arrange code to avoid hazards and stalls
 - Requires knowledge of the pipeline structure

Data Hazard Solution: Forwarding

- Forward result as soon as it is available
 - OK that it's not stored in RegFile yet

Arithmetic result
available in EX

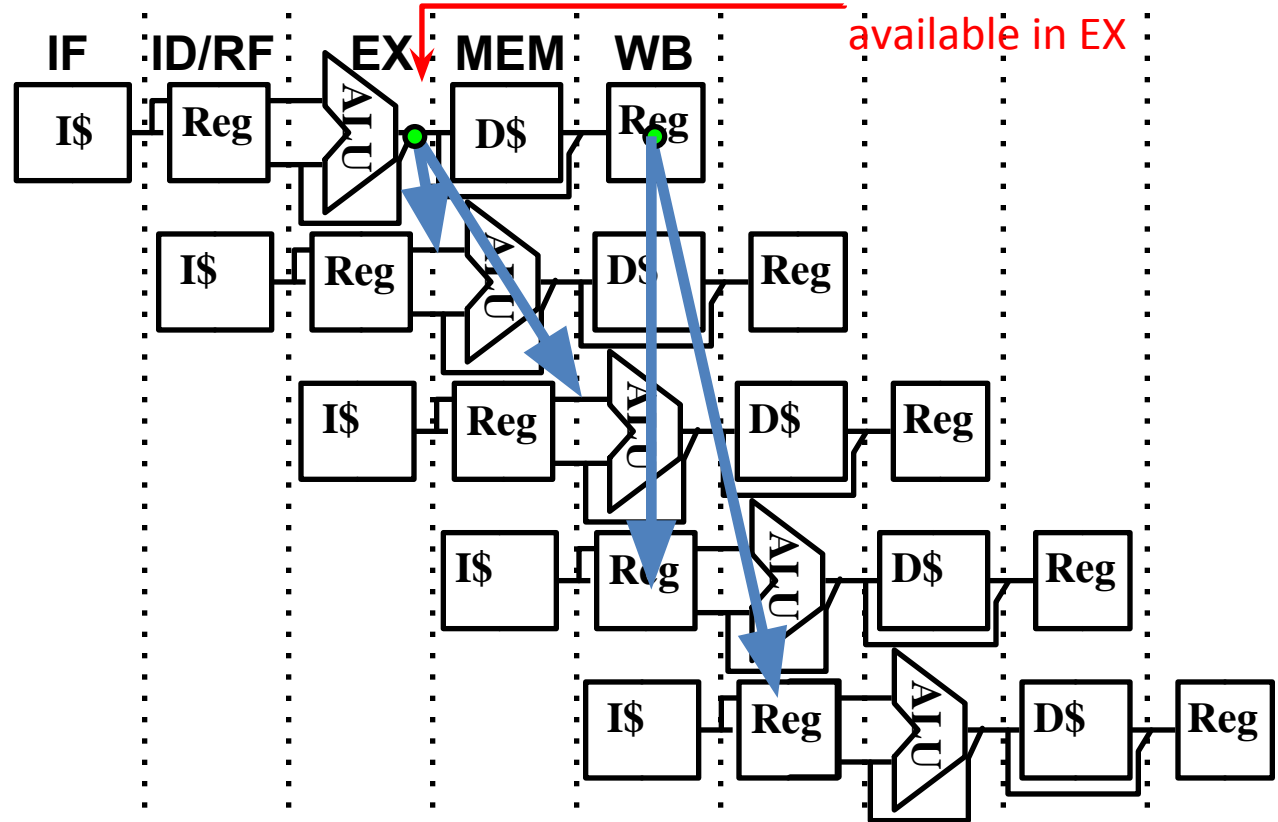
add t0, t1, t2

sub t4, t0, t3

and t5, t0, t6

or t7, t0, t8

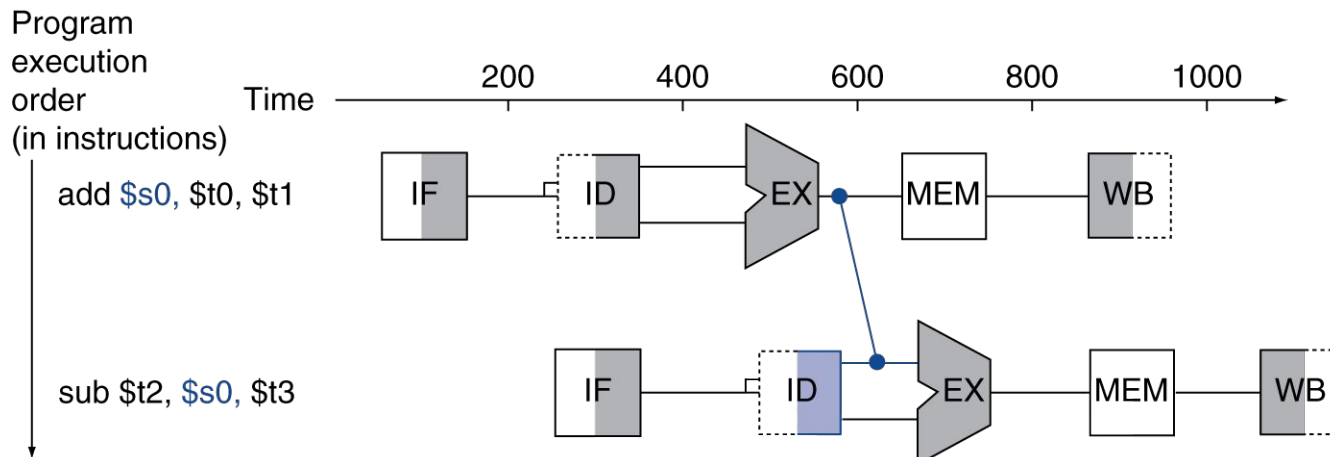
xor t9, t0, t10



Forwarding: grab operand from pipeline stage,
rather than register file

Forwarding (aka Bypassing)

- Use result when it is computed
 - Don't wait for it to be stored in a register
 - Requires extra connections in the datapath



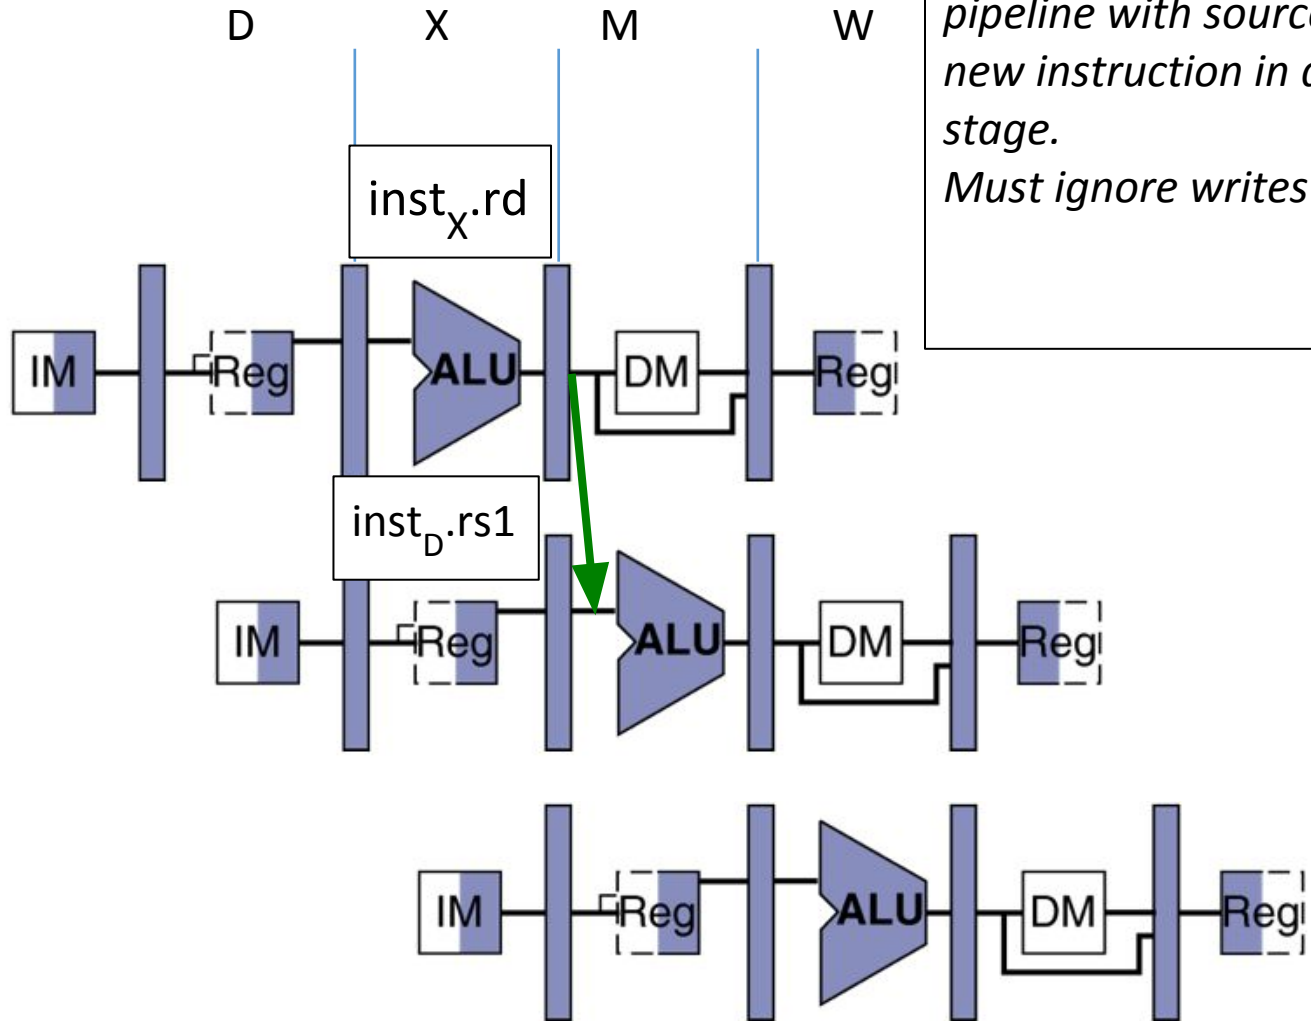
Detect Need for Forwarding (example)

*Compare destination of older instructions in pipeline with sources of new instruction in decode stage.
Must ignore writes to x0!*

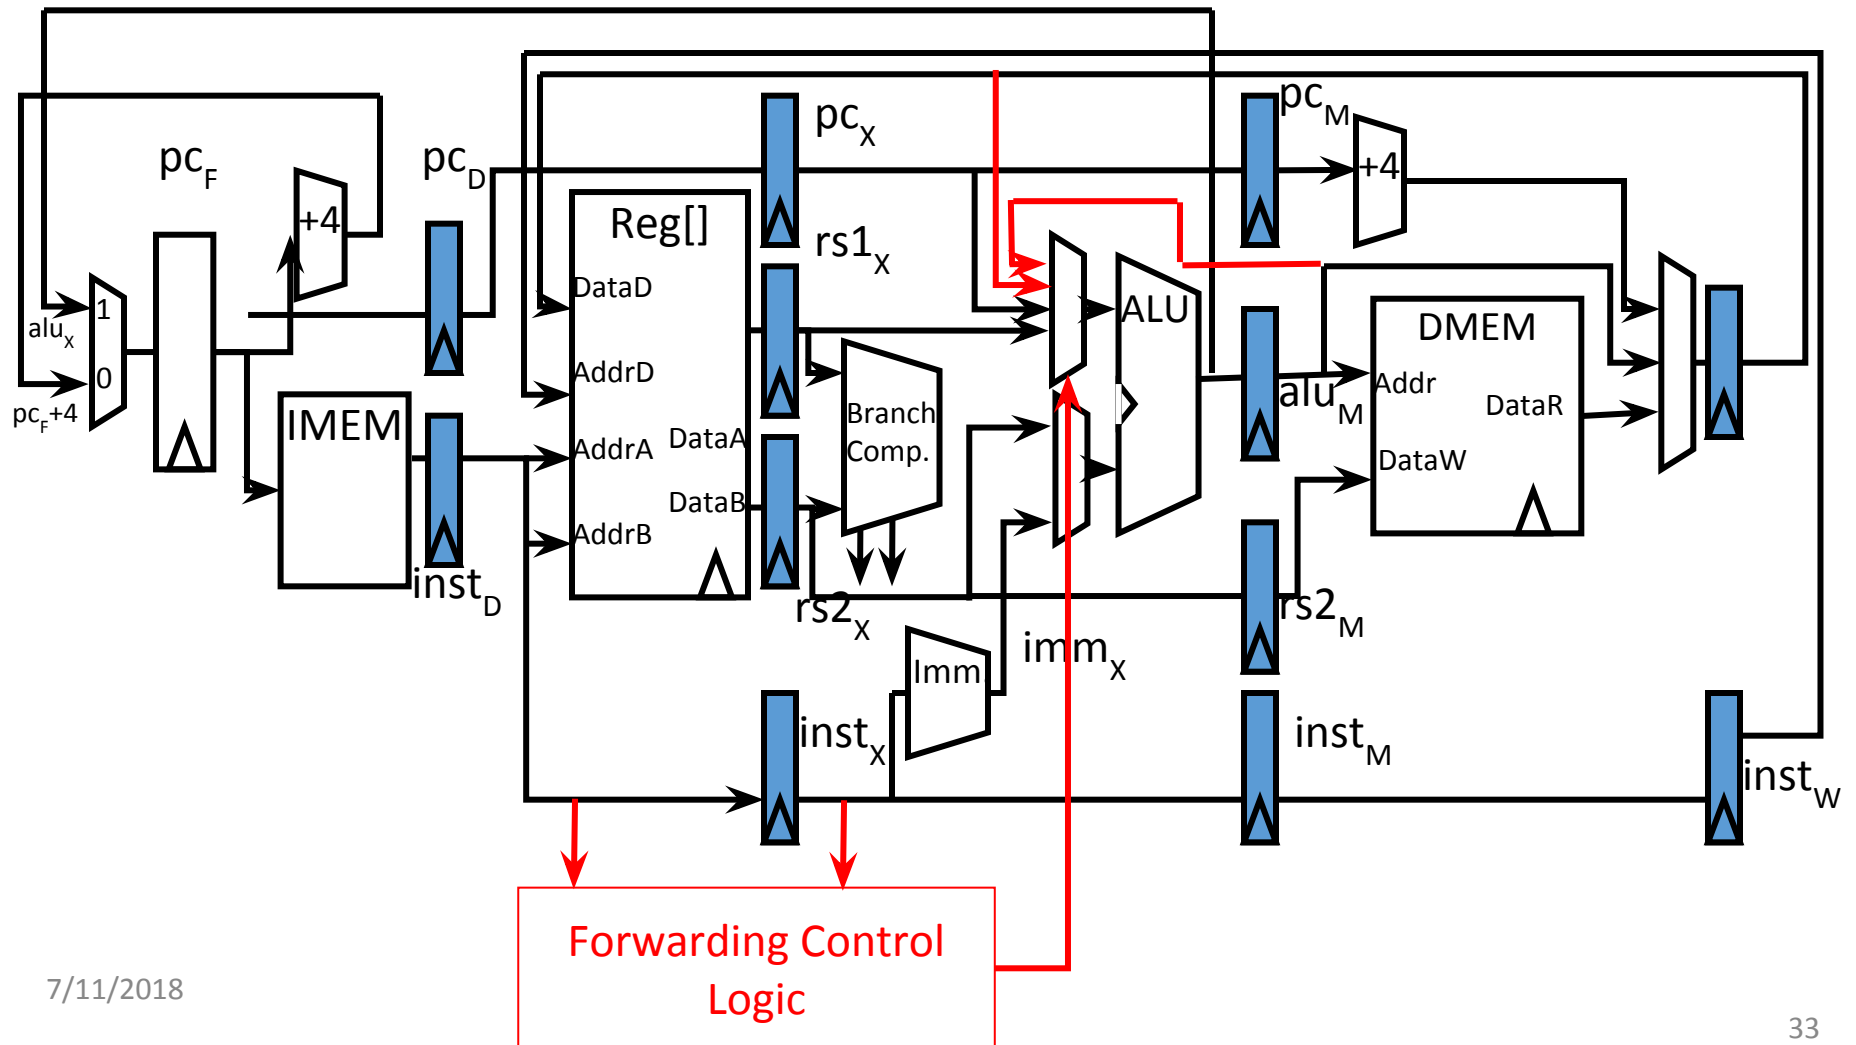
add t0, t1, t2

or t3, t0, t5

sub t6, t0, t3



Forwarding Path

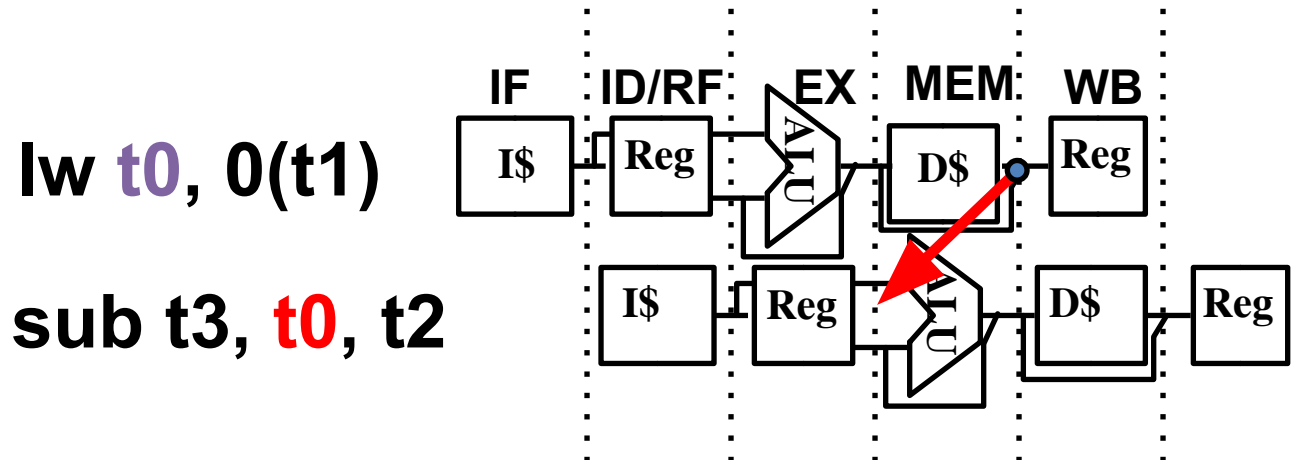


Agenda

- RISC-V Pipeline
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- Superscalar processors

Data Hazard: Loads (1/4)

- **Recall:** Dataflow backwards in time are hazards



- Can't solve all cases with forwarding
 - Must *stall* instruction dependent on load, then forward (more hardware)

Data Hazard: Loads (2/4)

- *Hardware stalls pipeline*
 - Called “hardware interlock”

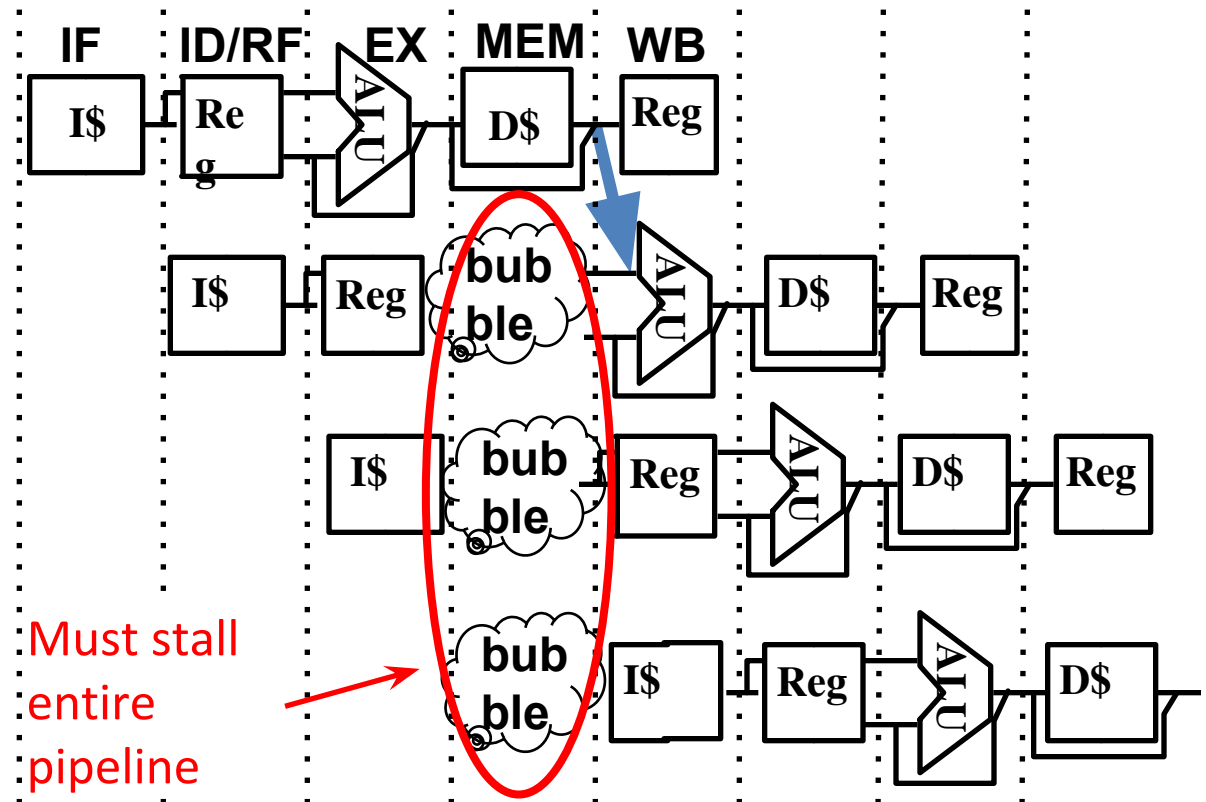
This is what happens in hardware in a “hardware interlock”

lw t0, 0(t1)

sub t3, t0, t2

and t5, t0, t4

or t7, t0, t6



Data Hazard: Loads (3/4)

- Stall is equivalent to `nop`

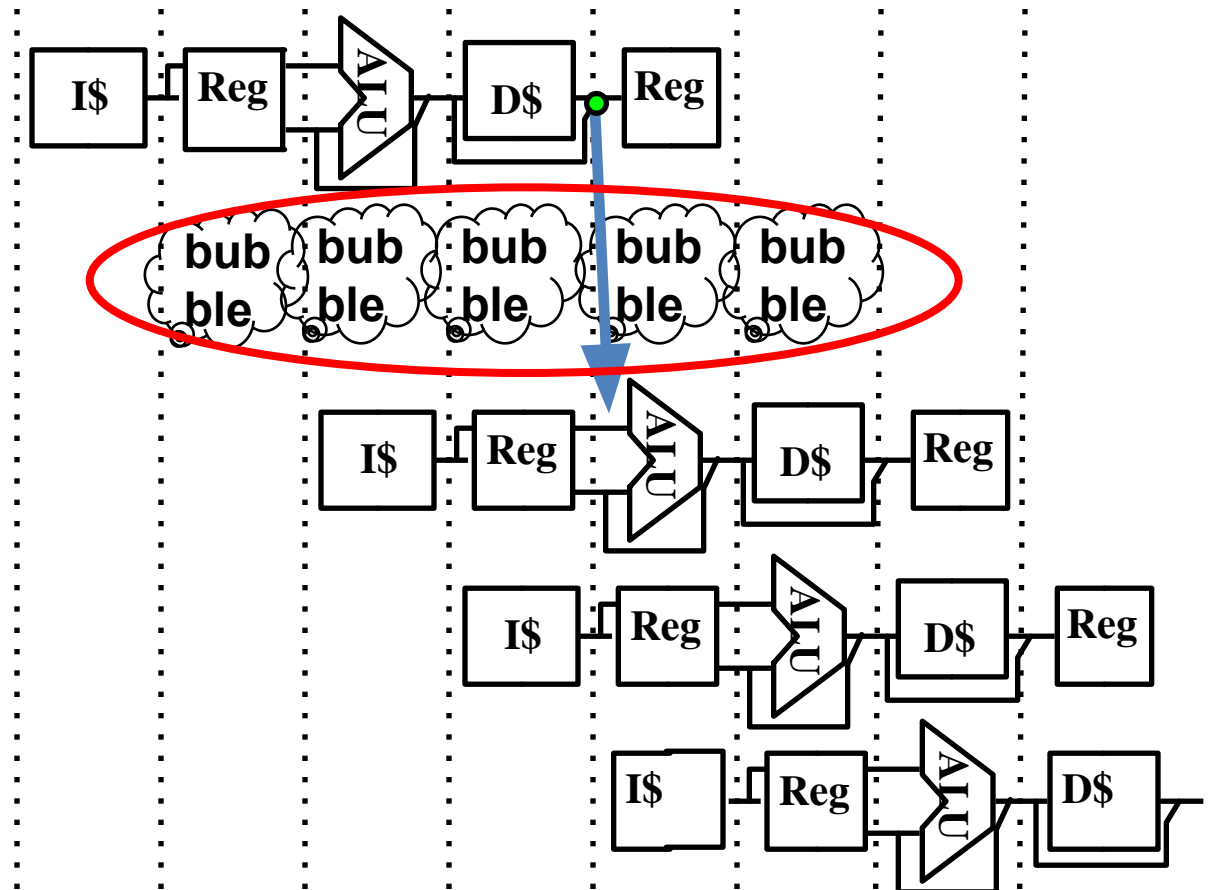
`lw t0, 0(t1)`

`nop`

`sub t3, t0, t2`

`and t5, t0, t4`

`or t7, t0, t6`



Data Hazard: Loads (4/4)

- Slot after a load is called a *load delay slot*
 - If that instruction uses the result of the load, then the hardware will stall for one cycle
 - Equivalent to inserting an explicit **nop** in the slot
 - except the latter uses more code space
 - Performance loss
- **Idea:** Let the compiler/assembler put an unrelated instruction in that slot → no stall!

Code Scheduling to Avoid Stalls

- Reorder code to avoid use of load result in the next instruction!
- RISC-V code for $D=A+B$; $E=A+C$;

Original Order:

```
lw t1, 0(t0)
lw t2, 4(t0)
add t3, t1, t2
sw t3, 12(t0)
lw t4, 8(t0)
add t5, t1, t4
sw t5, 16(t0)
```

Stall!

Stall!

13 cycles

Alternative:

```
lw t1, 0(t0)
lw t2, 4(t0)
lw t4, 8(t0)
add t3, t1, t2
sw t3, 12(t0)
add t5, t1, t4
sw t5, 16(t0)
```

11 cycles

Break!



Agenda

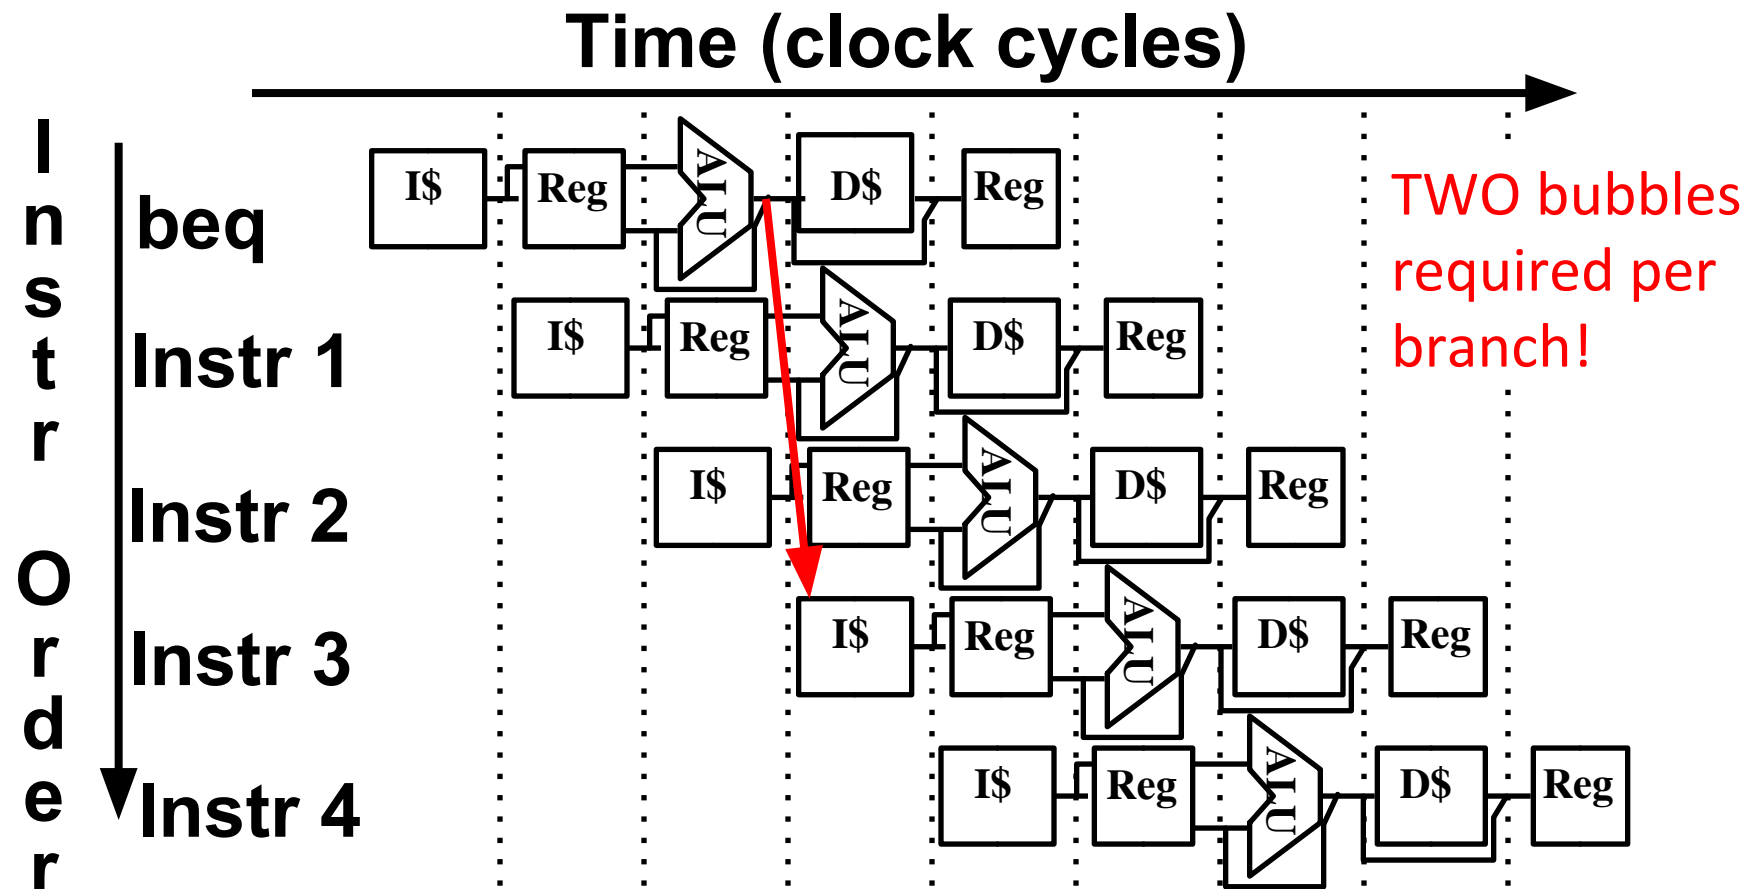
- RISC-V Pipeline
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - **Control**
- Superscalar processors

3. Control Hazards

- Branch (`beq`, `bne`) determines flow of control
 - Fetching next instruction depends on branch outcome
 - Pipeline can't always fetch correct instruction
 - Still working on ID stage of branch
- **Simple Solution:** Stall on *every* branch until we have the new PC value
 - How long must we stall?

Branch Stall

- When is comparison result available?

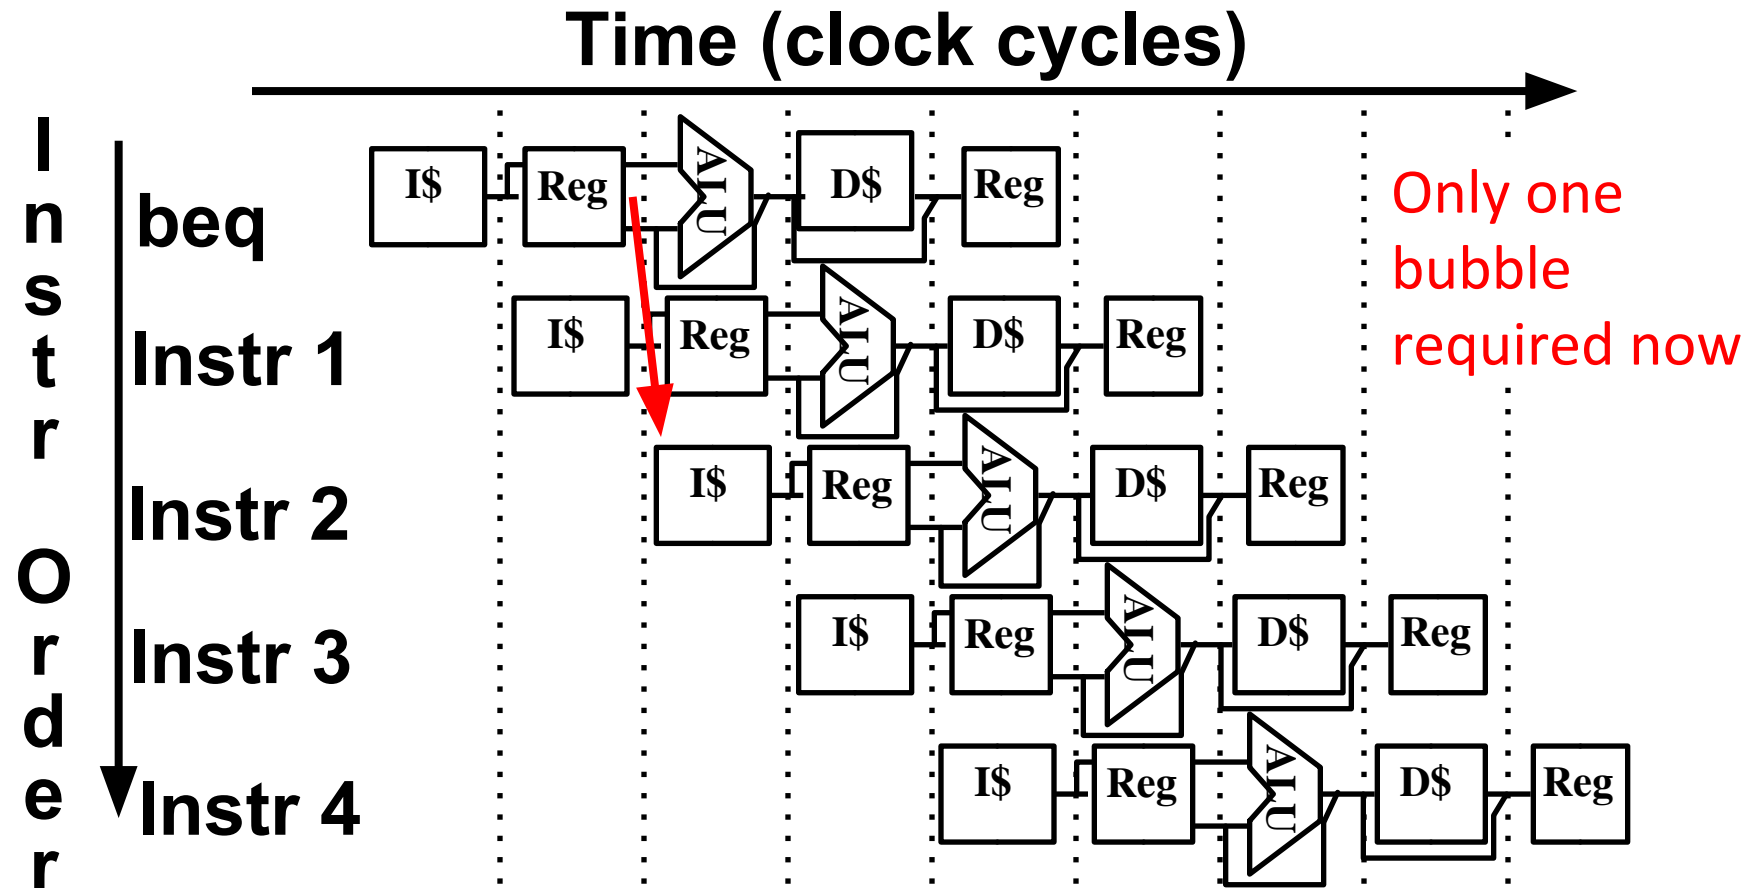


3. Control Hazard: Branching

- **Option #1:** Moving **branch comparator** to ID stage
 - As soon as instruction is decoded, immediately make a decision and set the new value of the PC
 - **Benefit:** Branch decision made in 2nd stage, so only one `nop` is needed instead of two
 - **Side Note:** Have to compute new PC value in ID instead of EX
 - Adds extra hardware and reduces redundancy
 - Branches are idle in EX, MEM, and WB

Improved Branch Stall

- When is comparison result available?

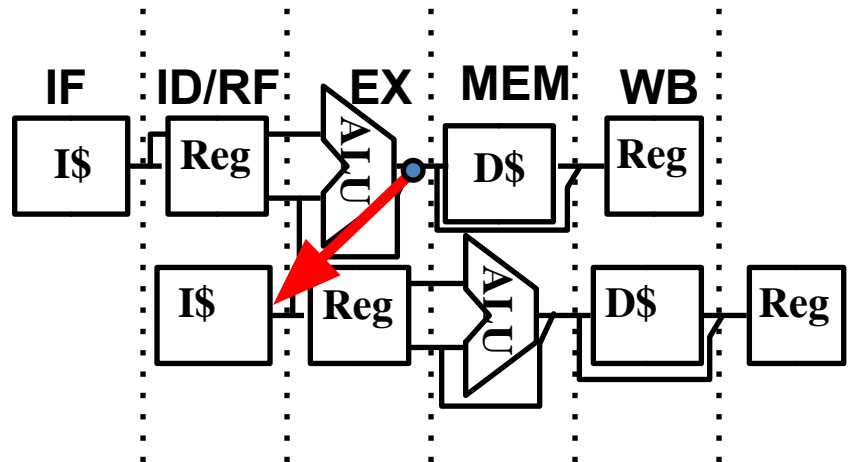


Data Hazard: Branches!

- **Recall:** Dataflow backwards in time are hazards

add **t0**, t0, t1

beq x0, **t0**, foo



- Now that **t0** is needed earlier (ID instead of EX), we can't forward it to the beq's ID stage
 - Must *stall* after add, then forward (more hardware)

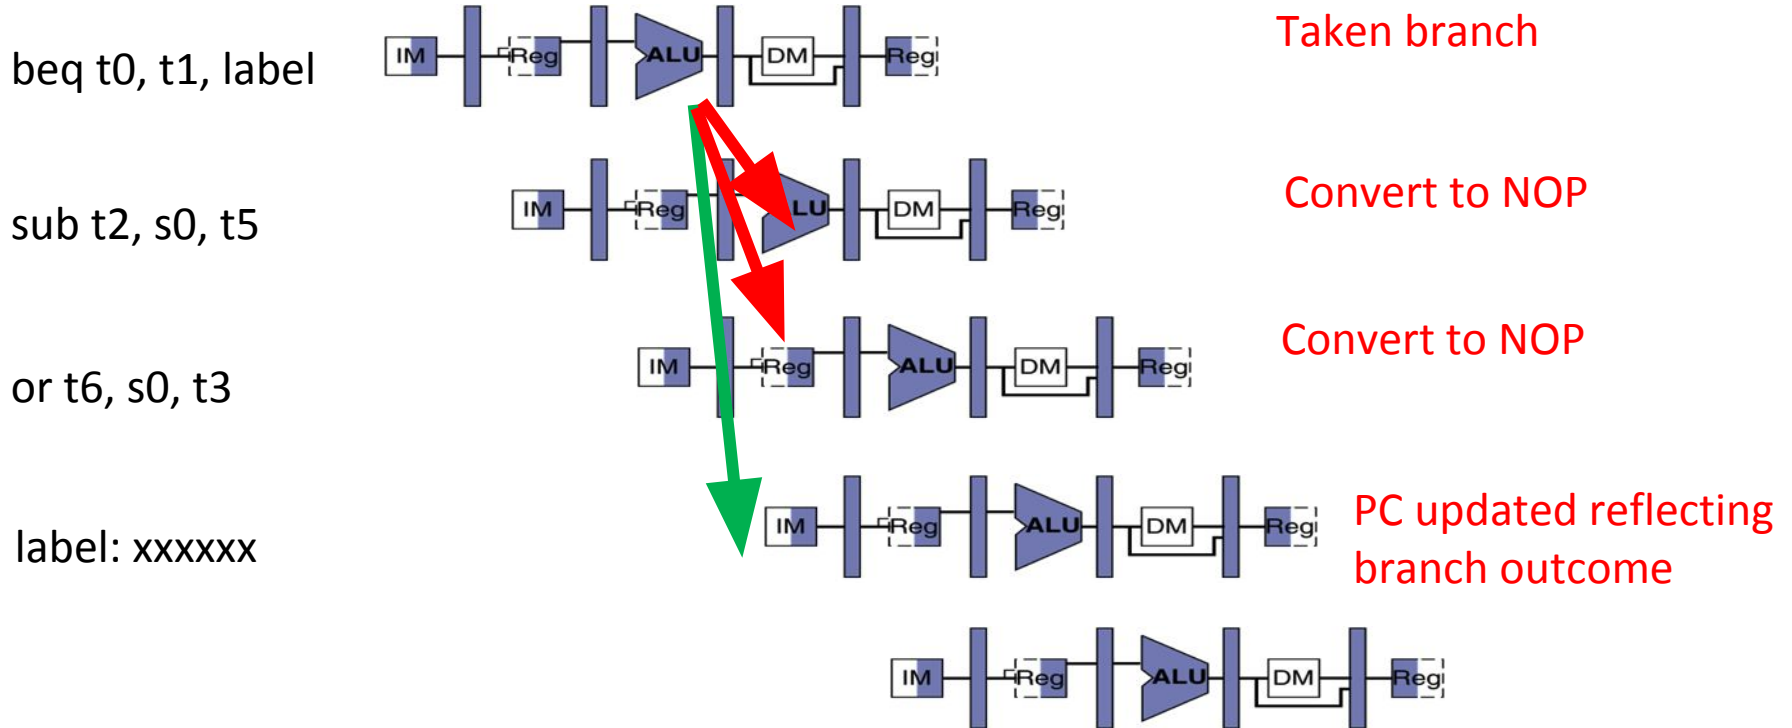
Observations

- **Takeaway:** Moving **branch comparator** to ID stage would add extra hardware, reduce redundancy, and introduce new problems
- Can we work with the nature of branches?
 - If branch not taken, then instructions fetched sequentially after branch are correct
 - If branch or jump taken, then need to flush incorrect instructions from pipeline by converting to NOPs

3. Control Hazard: Branching

- **RISC-V Solution:** *Branch Prediction* – guess outcome of a branch, fix afterwards if necessary
 - Must cancel (*flush*) all instructions in pipeline that depended on guess that was wrong
 - How many instructions do we end up flushing?

Kill Instructions after Branch if Taken



Branch Prediction


beq t0, t1, label

label:

.....

....

.



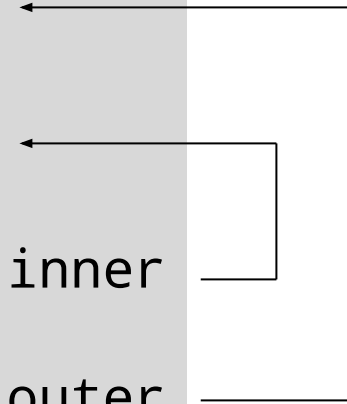
Dynamic Branch Prediction

- Branch penalty is more significant in deeper pipelines
- Use *dynamic branch prediction*
 - Have branch prediction buffer (a.k.a. branch history table) that stores outcomes (taken/not taken) indexed by recent branch instruction addresses
 - To execute a branch
 - Check table and predict the same outcome for next fetch
 - If wrong, flush pipeline and flip prediction

1-Bit Predictor: Shortcoming

- Examine the code below, assuming both loops will be executed multiple times:

```
outer: ...  
      ...  
inner: ...  
      ...  
      beq ..., ..., inner  
      ...  
      beq ..., ..., outer
```



- Inner loop branches are predicted wrong twice!
 - Predict as taken on last iteration of inner loop
 - Then predict as not taken on first iteration of inner loop next time around

Agenda

- Structural Hazards
- Data Hazards
 - Forwarding
- Administtrivia
- Data Hazards (Continued)
 - Load Delay Slot
- **Control Hazards**
 - Branch and Jump Delay Slots
 - **Branch Prediction**

Question: For each code sequences below, choose one of the statements below:

1:

```
lw $t0, 0($t0)
add $t1, $t0, $t0
```

2:

```
add $t1, $t0, $t0
addi $t2, $t0, 5
addi $t4, $t1, 5
```

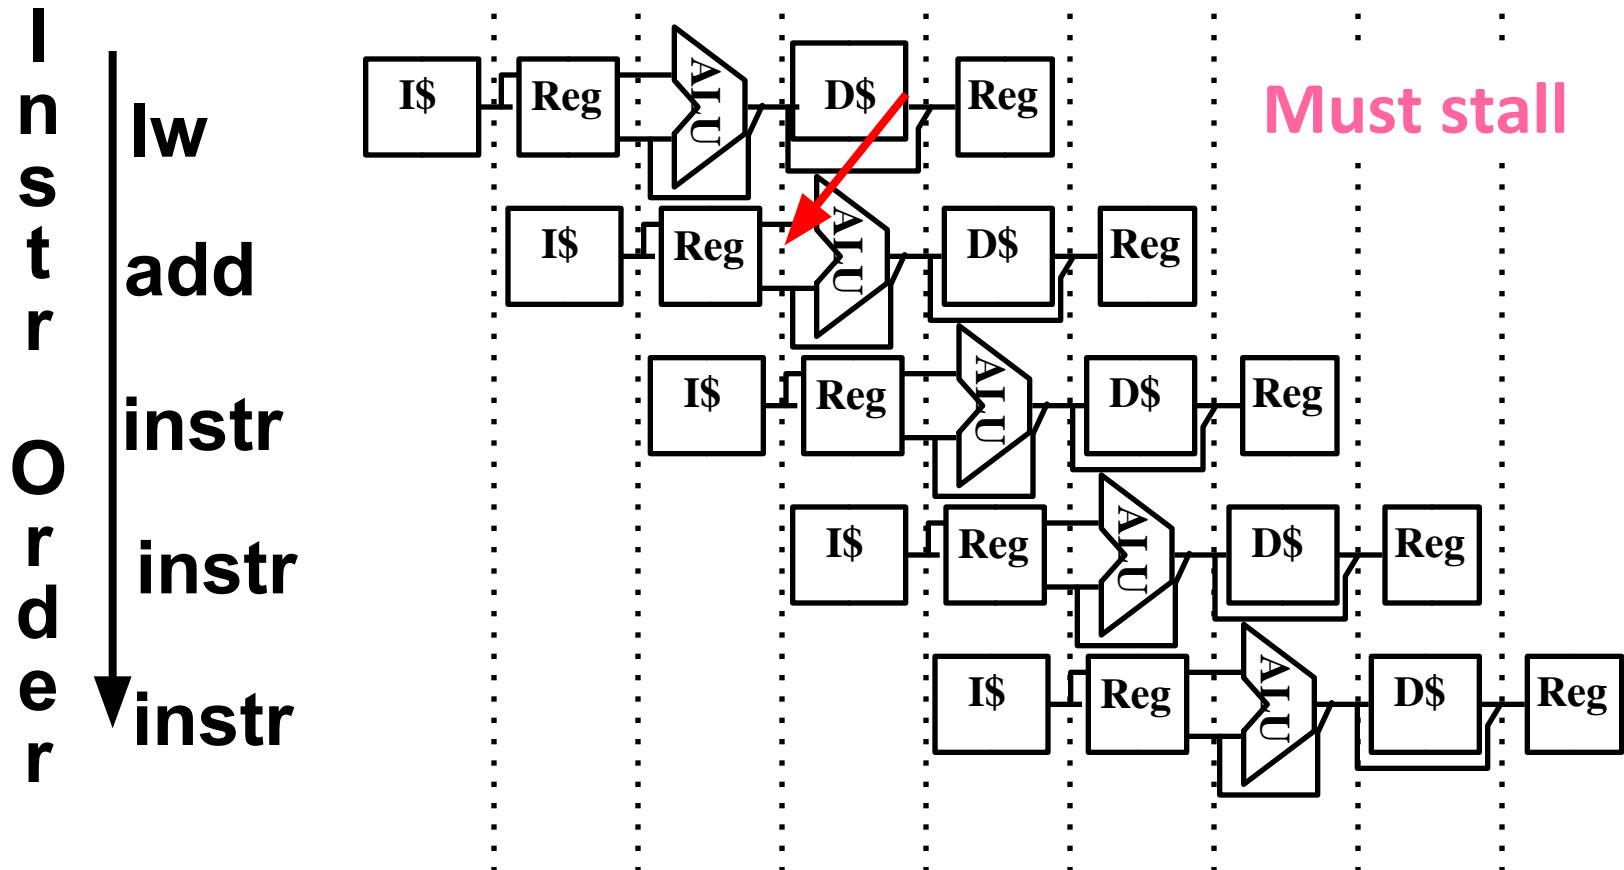
3:

```
addi $t1, $t0, 1
addi $t2, $t0, 2
addi $t3, $t0, 2
addi $t3, $t0, 4
addi $t5, $t1, 5
```

- A **No stalls as is**
- B **No stalls with forwarding**
- C **Must stall**

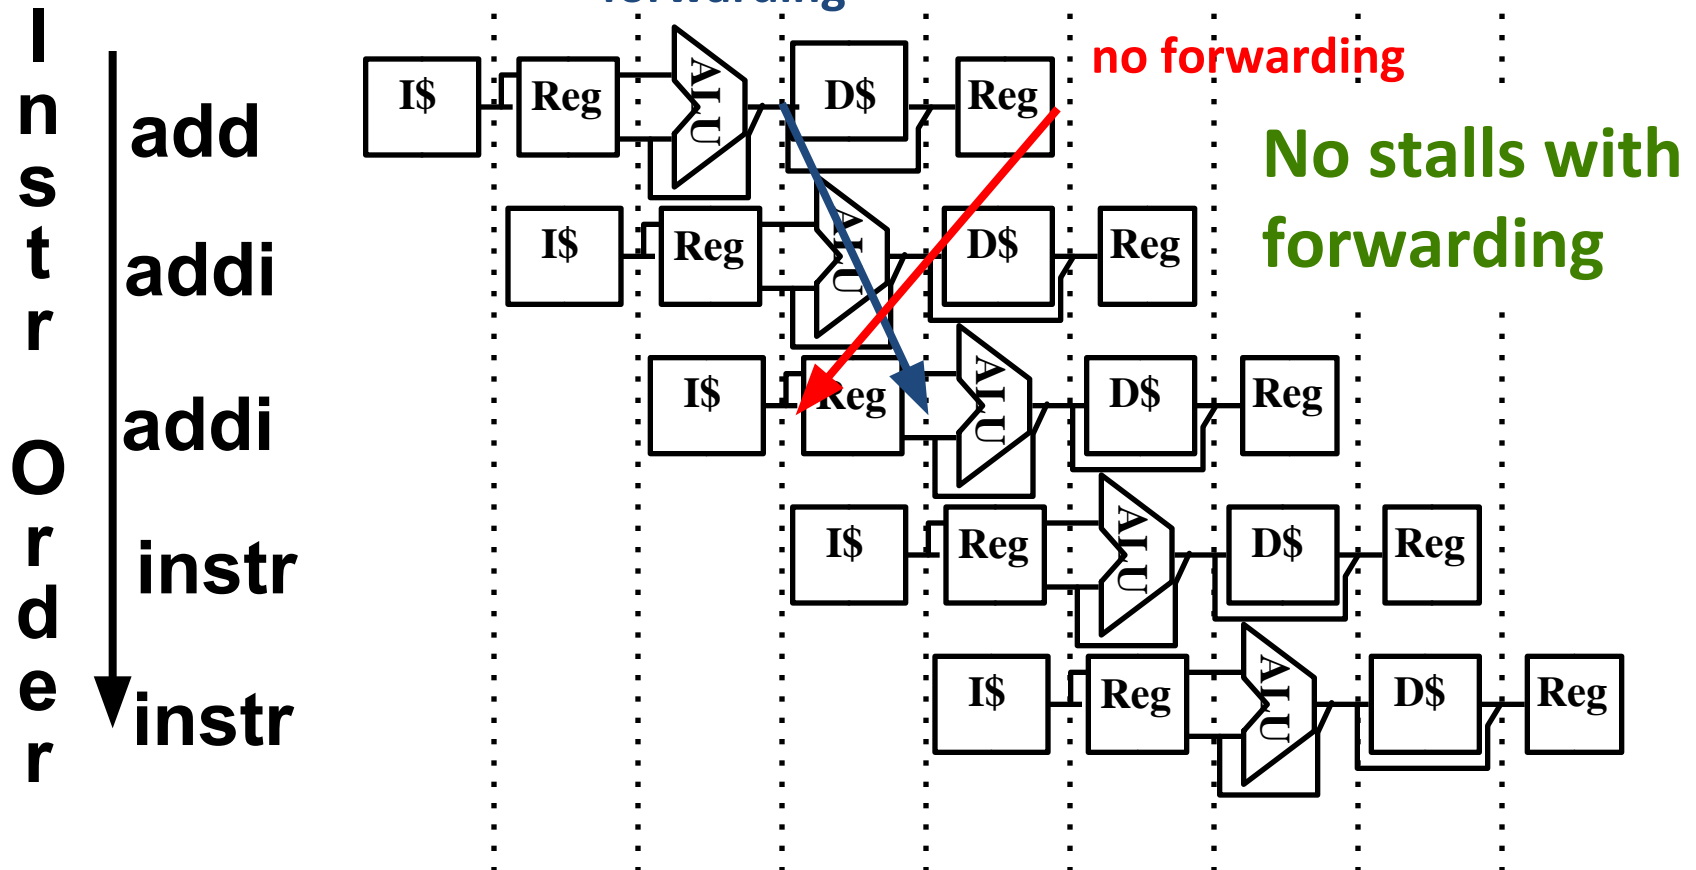
Code Sequence 1

Time (clock cycles)



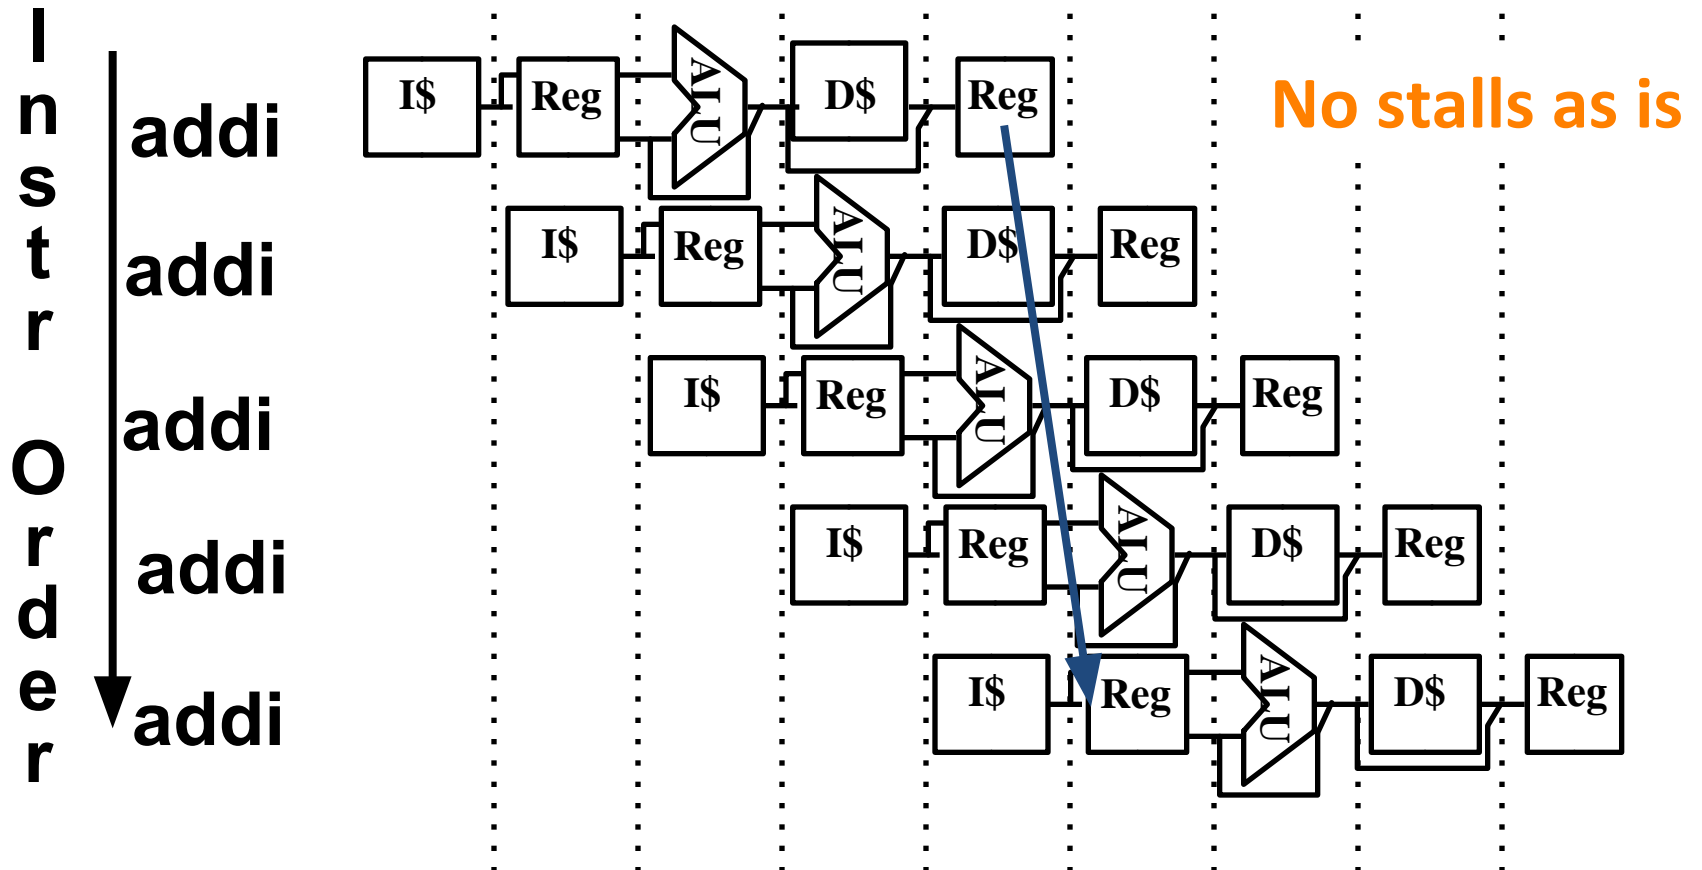
Code Sequence 2

Time (clock cycles)



Code Sequence 3

Time (clock cycles)



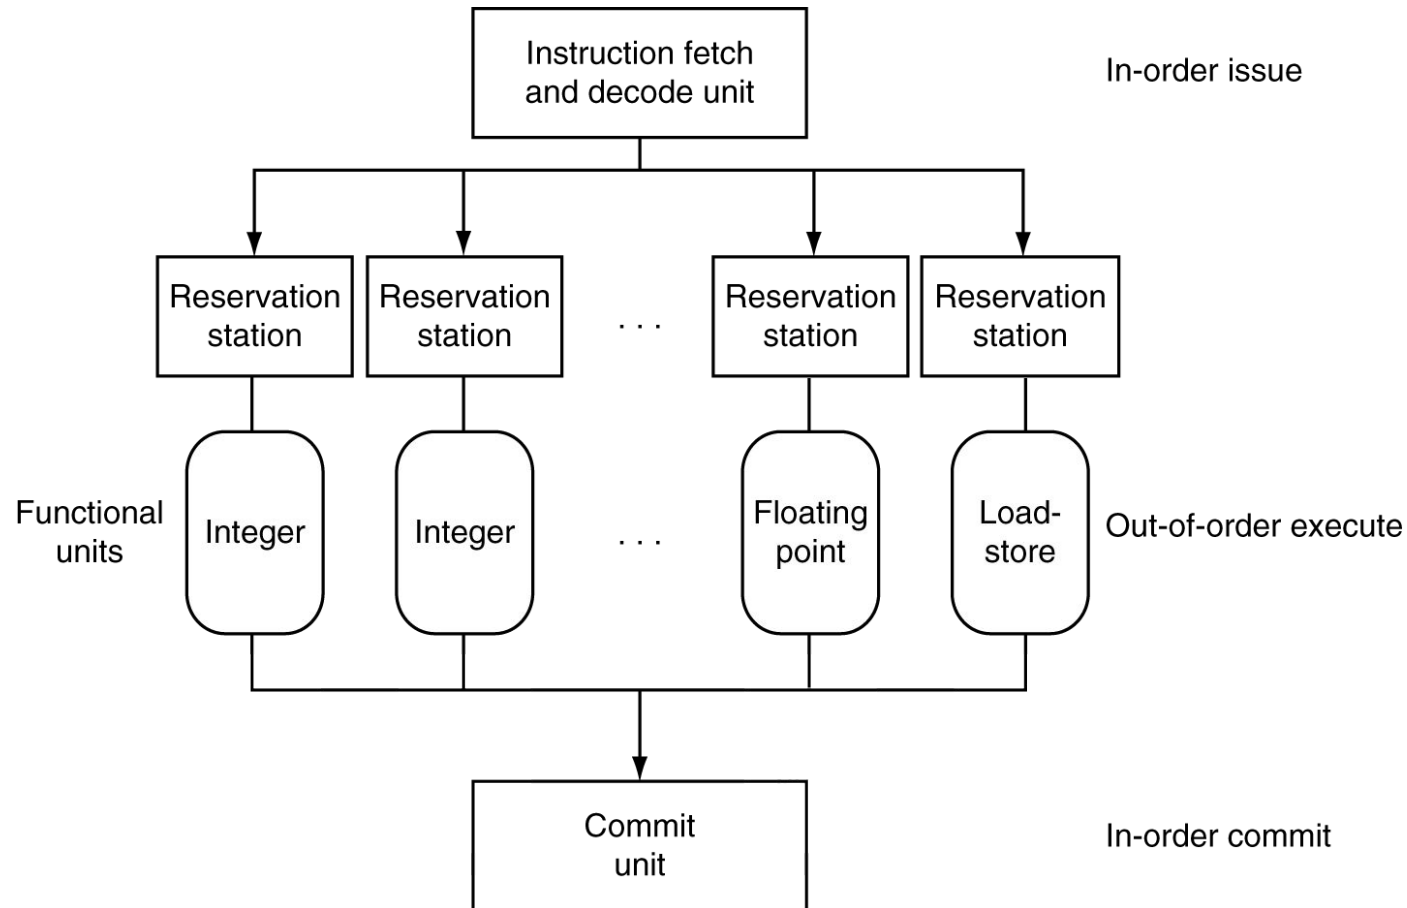
Agenda

- RISC-V Pipeline
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- Superscalar processors

Increasing Processor Performance


1. Clock rate
 - Limited by technology and power dissipation
2. Pipelining
 - “Overlap” instruction execution
 - Deeper pipeline: 5 \Rightarrow 10 \Rightarrow 15 stages
 - Less work per stage \rightarrow shorter clock cycle
 - But more potential for hazards ($CPI > 1$)
3. Multi-issue “super-scalar” processor
 - Multiple execution units (ALUs)
 - Several instructions executed simultaneously
 - $CPI < 1$ (ideally)

Superscalar Processor



P&H p. 340

Benchmark: CPI of Intel Core i7



CPI of Intel Core i7 920 running SPEC2006 integer benchmarks.

P&H p. 350

Summary

- Hazards reduce effectiveness of pipelining
 - Cause stalls/bubbles
- Structural Hazards
 - Conflict in use of a datapath component
- Data Hazards
 - Need to wait for result of a previous instruction
- Control Hazards
 - Address of next instruction uncertain/unknown
- Superscalar processors use multiple execution units for additional instruction level parallelism
 - Performance benefit highly code dependent

Extra Slides

Pipelining and ISA Design

- RISC-V ISA designed for pipelining
 - All instructions are 32-bits
 - Easy to fetch and decode in one cycle
 - Versus x86: 1- to 15-byte instructions
 - Few and regular instruction formats
 - Decode and read registers in one step
 - Load/store addressing
 - Calculate address in 3rd stage, access memory in 4th stage
 - Alignment of memory operands
 - Memory access takes only one cycle

Superscalar Processor

- Multiple issue “superscalar”
 - Replicate pipeline stages \Rightarrow multiple pipelines
 - Start multiple instructions per clock cycle
 - $CPI < 1$, so use Instructions Per Cycle (IPC)
 - E.g., 4GHz 4-way multiple-issue
 - 16 BIPS, peak $CPI = 0.25$, peak $IPC = 4$
 - Dependencies reduce this in practice
- “Out-of-Order” execution
 - Reorder instructions dynamically in hardware to reduce impact of hazards
- *CS152 discusses these techniques!*